

面向RISC-V异构AI芯片的 “大编译器”设计和实现



伍华林

兆松科技（武汉）有限公司

目录

CONTENTS

01 | 动机

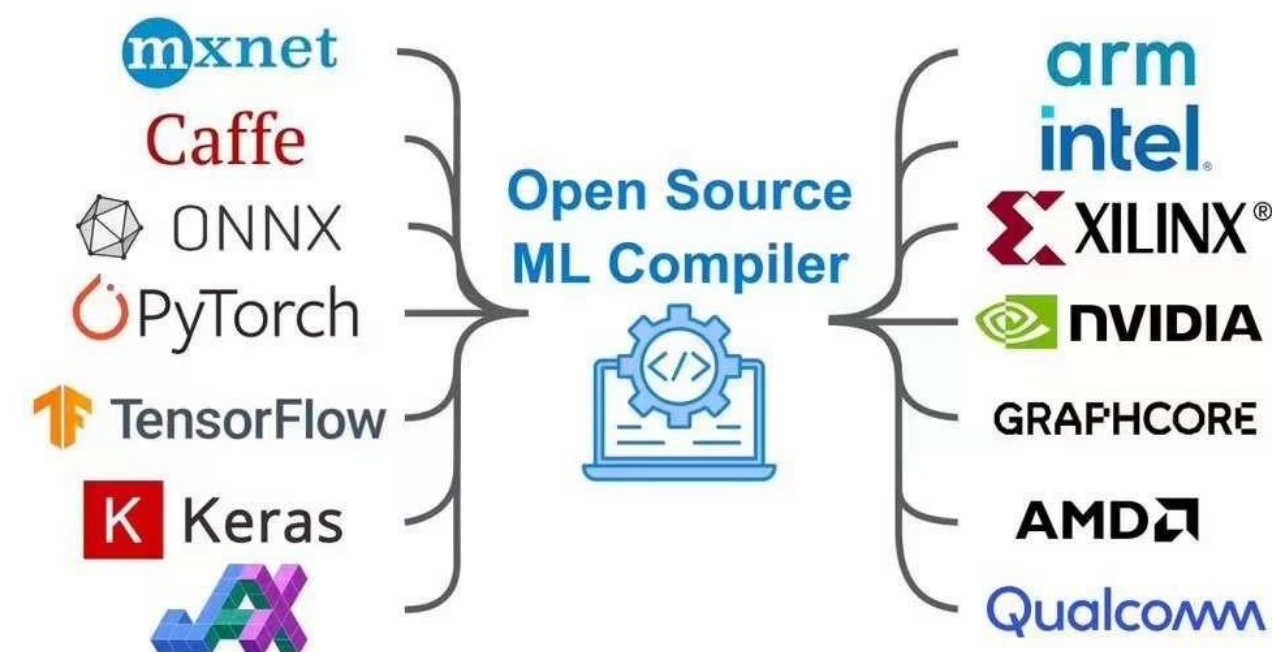
02 | “大编译器” 设计

03 | “大编译器” 实现


04 | 总结

- 软件1.0时代 -> 软件2.0时代
- 新模型, AI加速器层出不穷
- 模型 x 框架 x 加速器: 指数级增长
- 算子库维护成本, CUDA兼容问题

Expectation



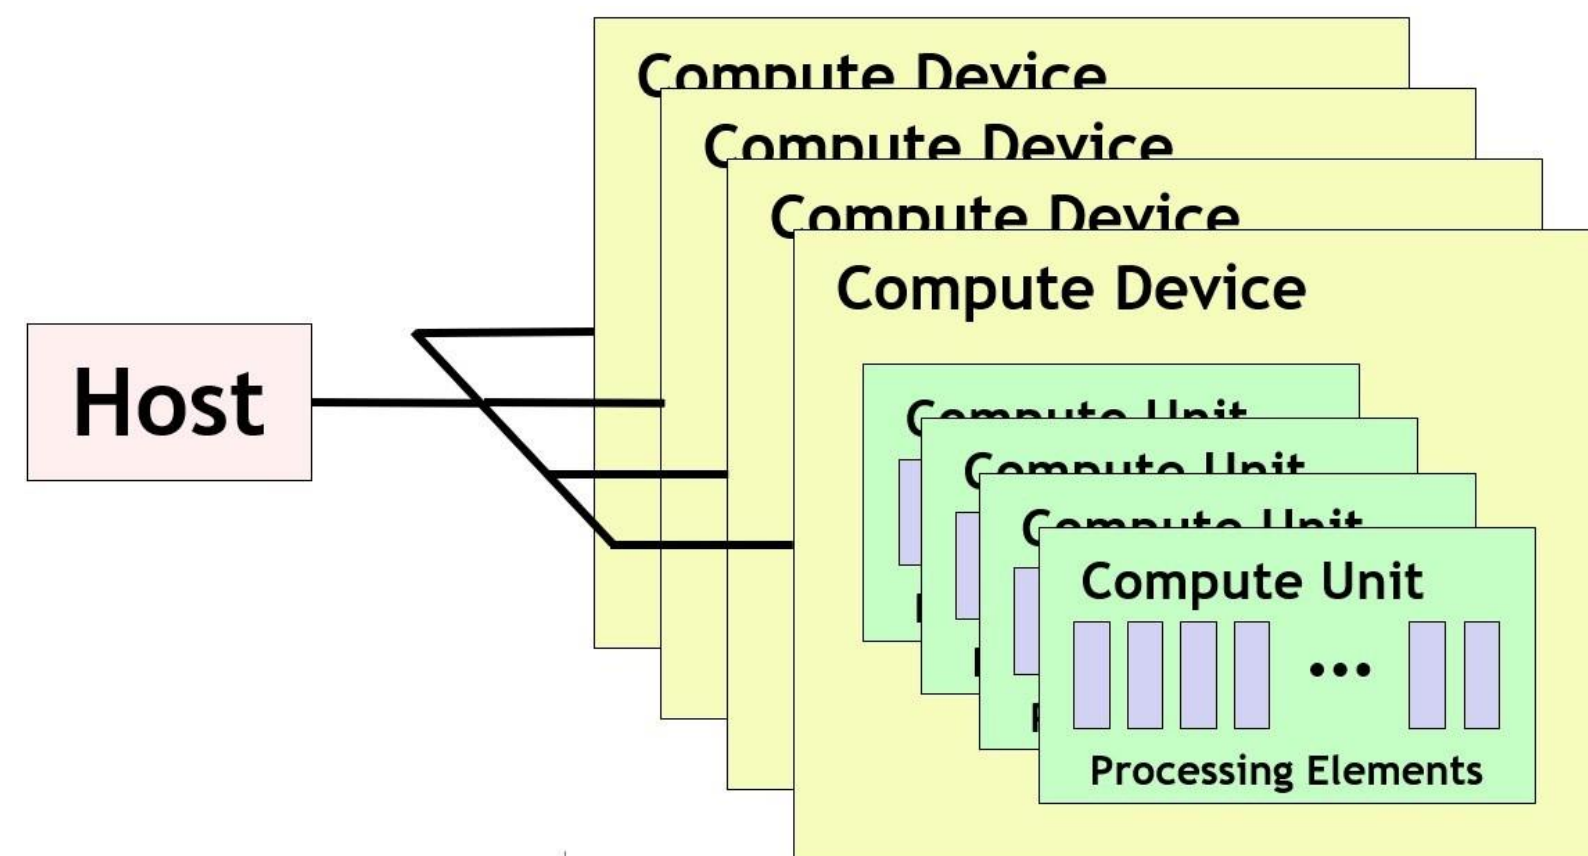
Reality



- LLVM对加速器的支持成熟且容易
- LLVM擅长做标量优化
- MLIR提供了多层抽象，适合做LLVM不擅长的优化
- MLIR后端直接生成LLVM IR方言

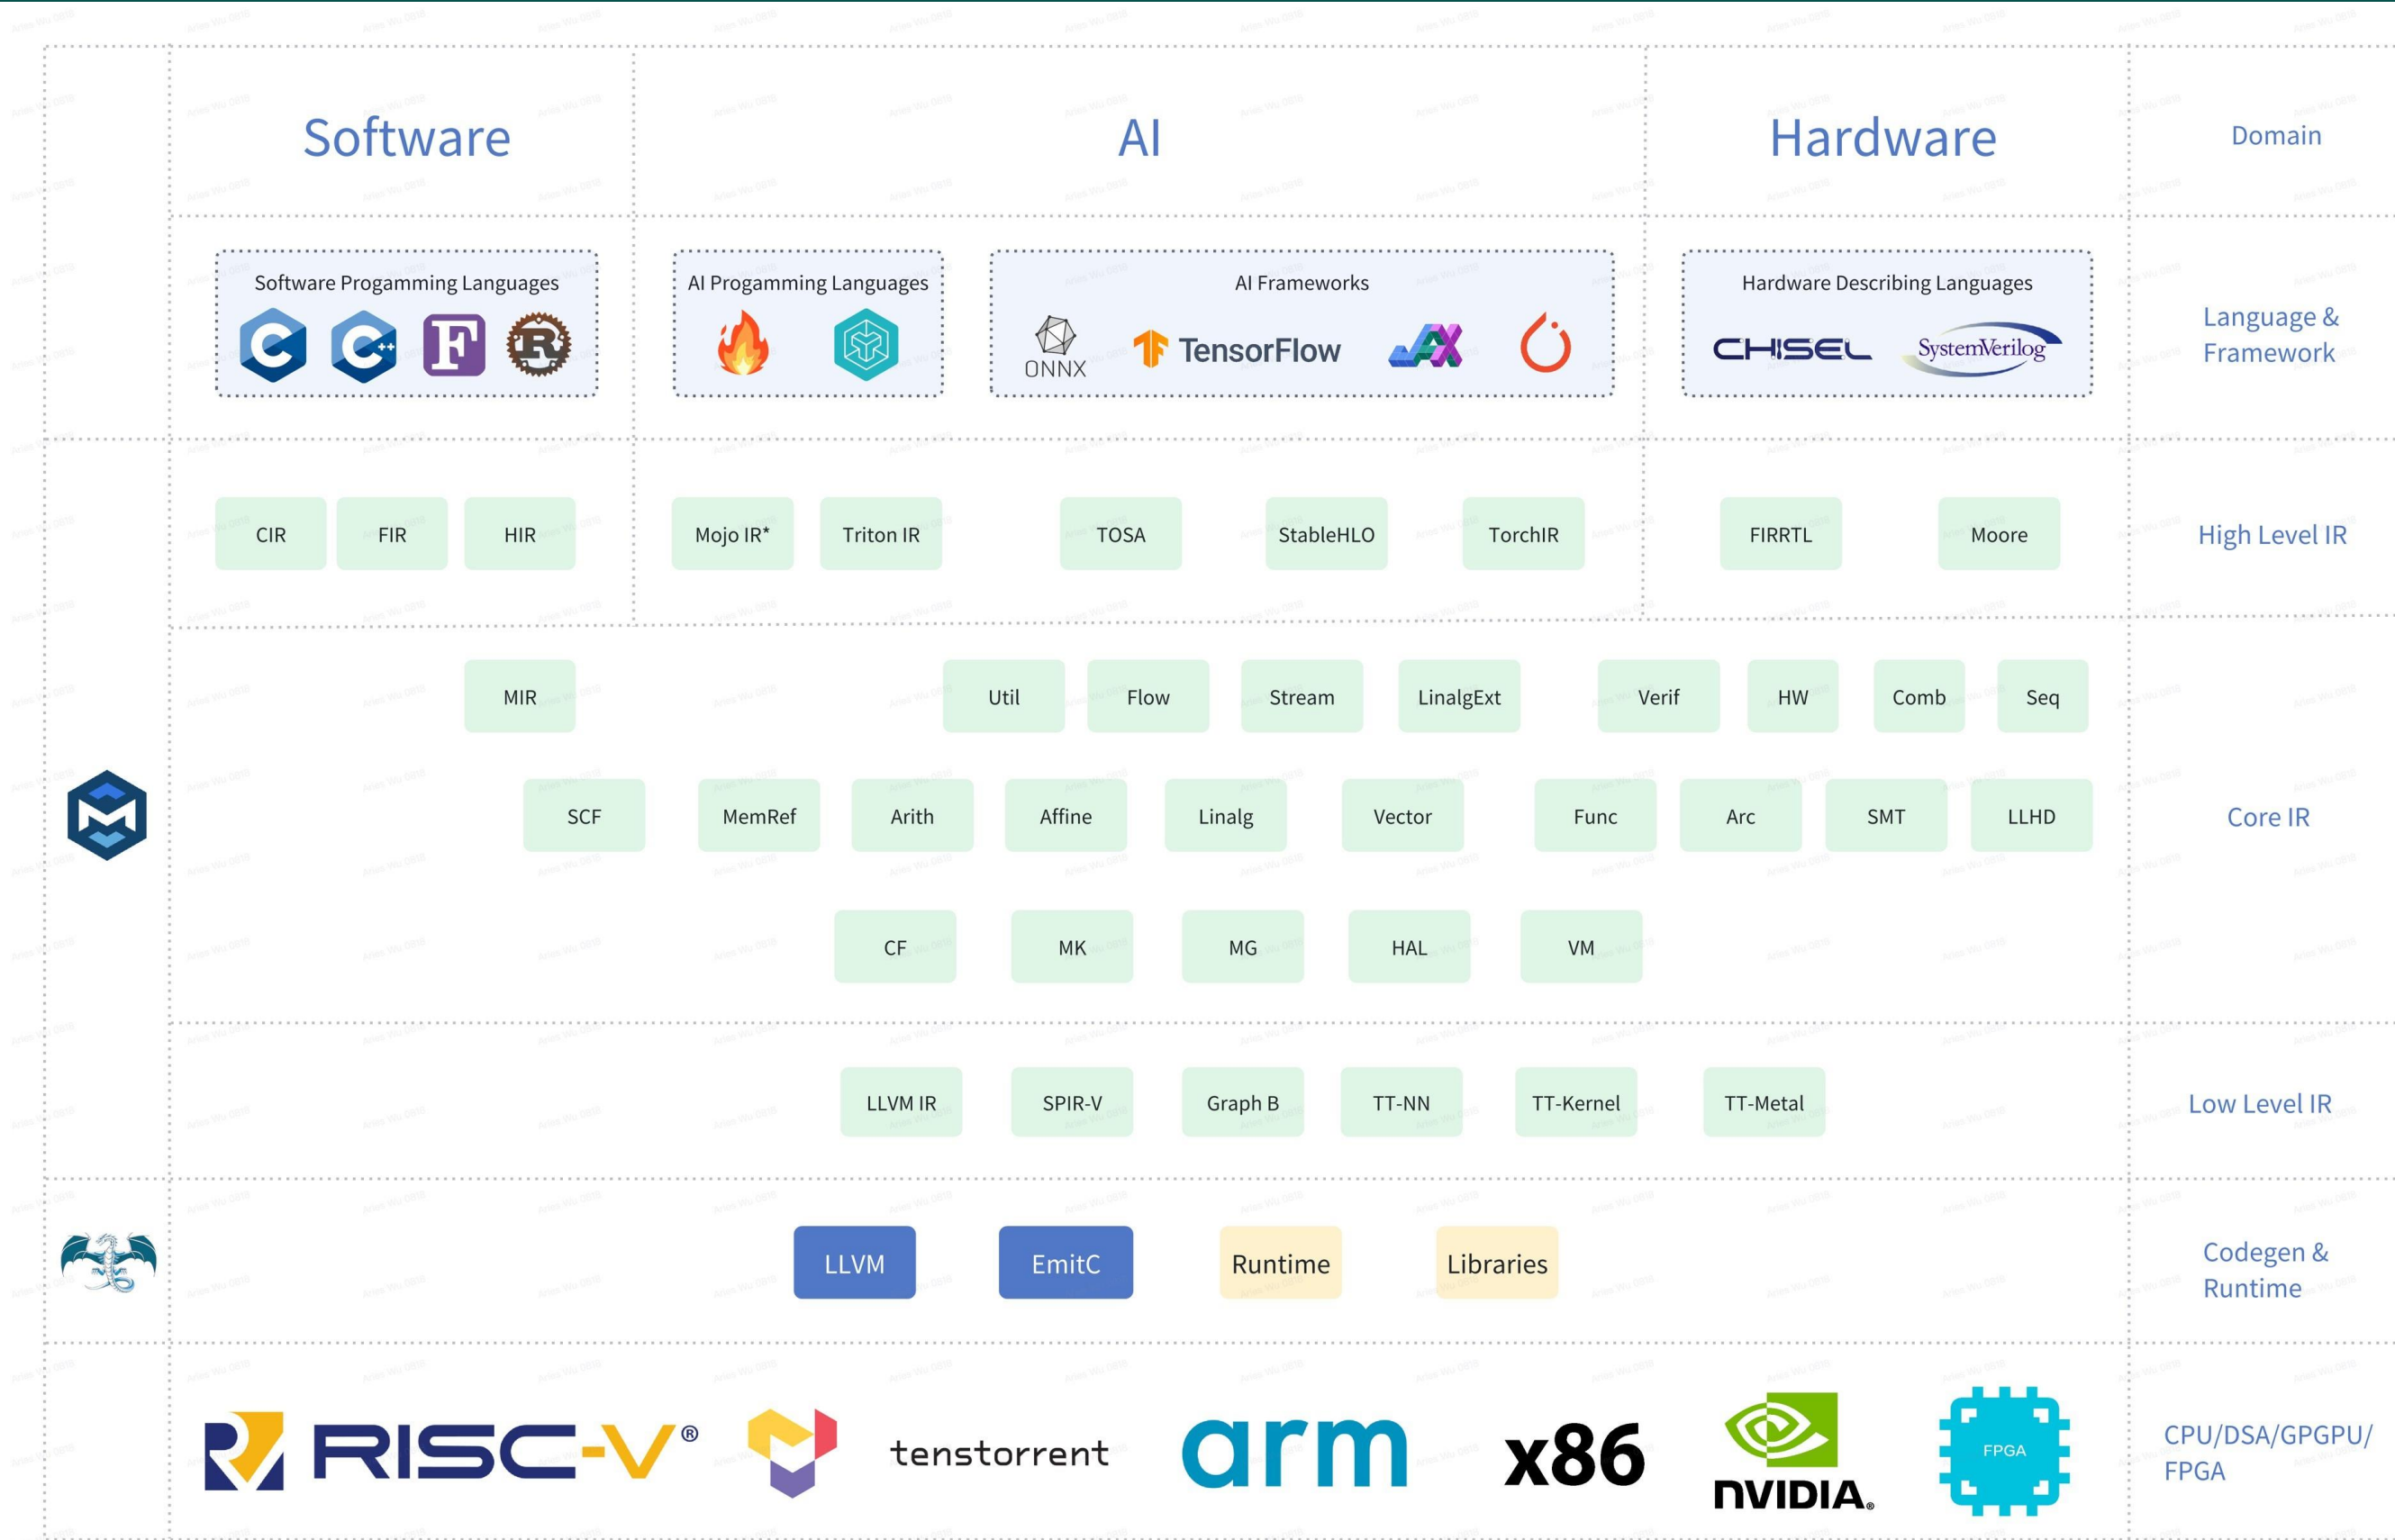


- 传统编译器：针对某一款特定指令集的芯片的编译器
- 大编译器
 - 对控制器和加速器的计算和任务程序，以及两者之间的通信进行抽象
 - 自动生成异构计算芯片的所有代码（包括驱动，加速器代码）



异构计算系统案例 - OpenCL Platform Model

“大编译器” 设计



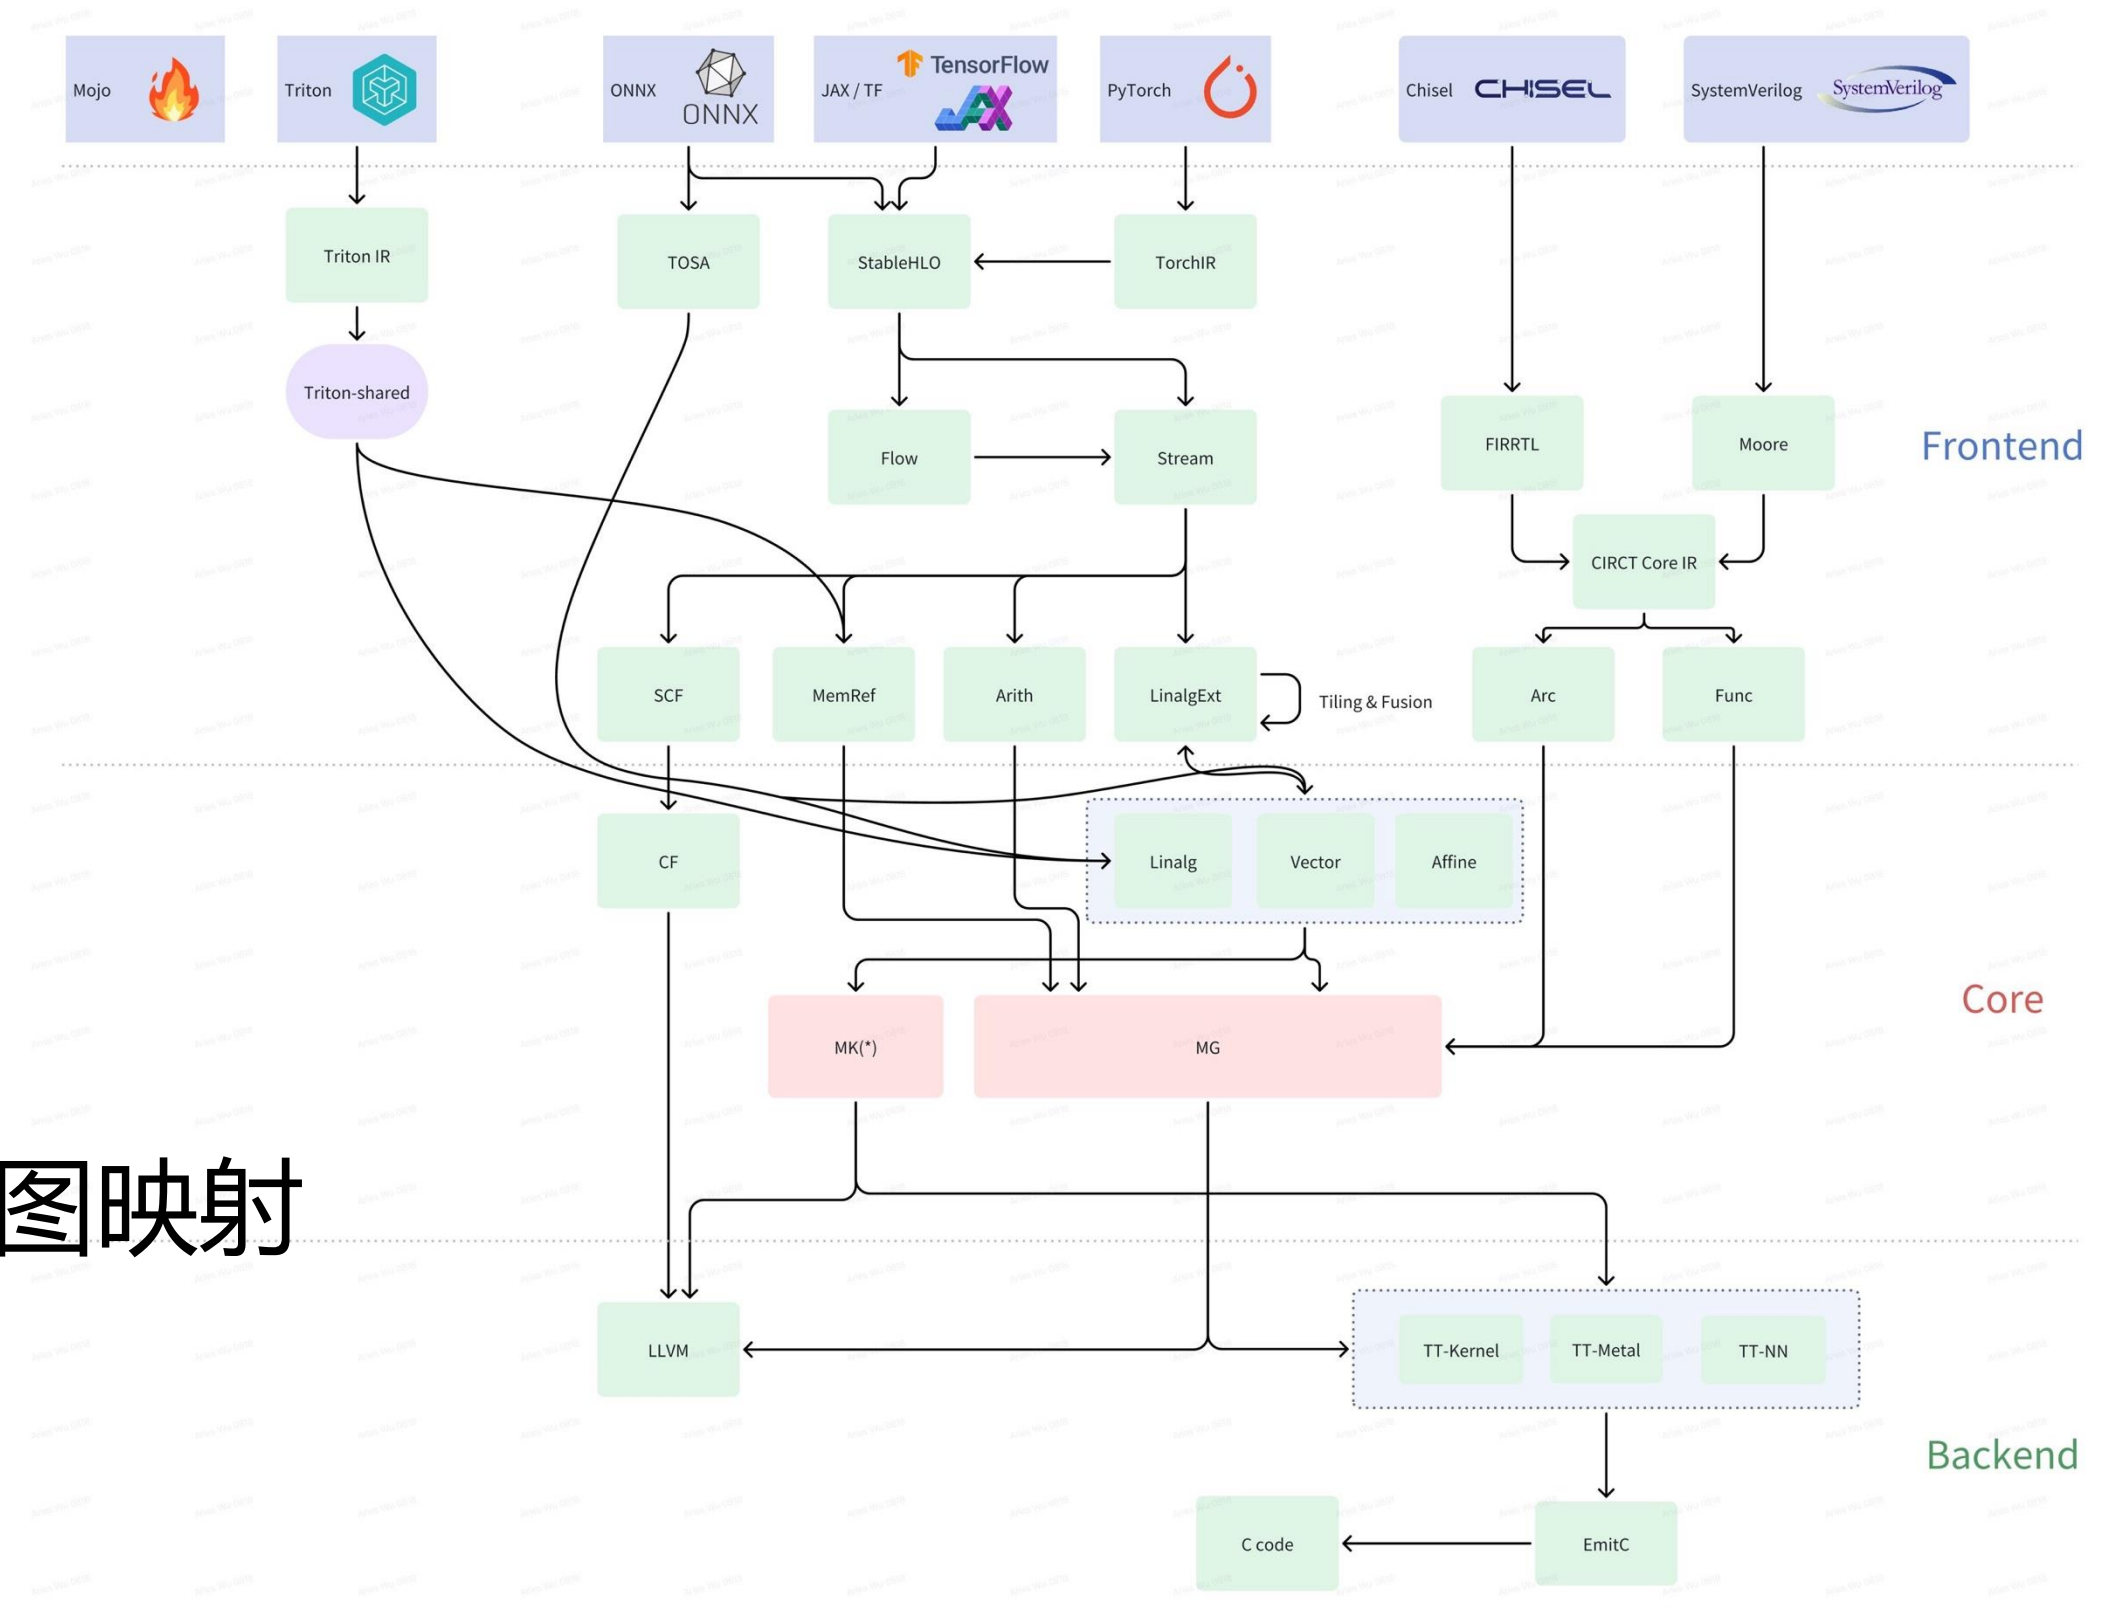
“大编译器” 设计 – MLIR前端



- 基于IREE/MLIR的一些核心方言
- 支持Triton/Mojo*编程语言
- 支持多种模型文件导入
- 支持Dataflow芯片计算图生成
- 统一的算子和计算图抽象层
- 支持RTL计算图到Dataflow计算图映射

*未来支持

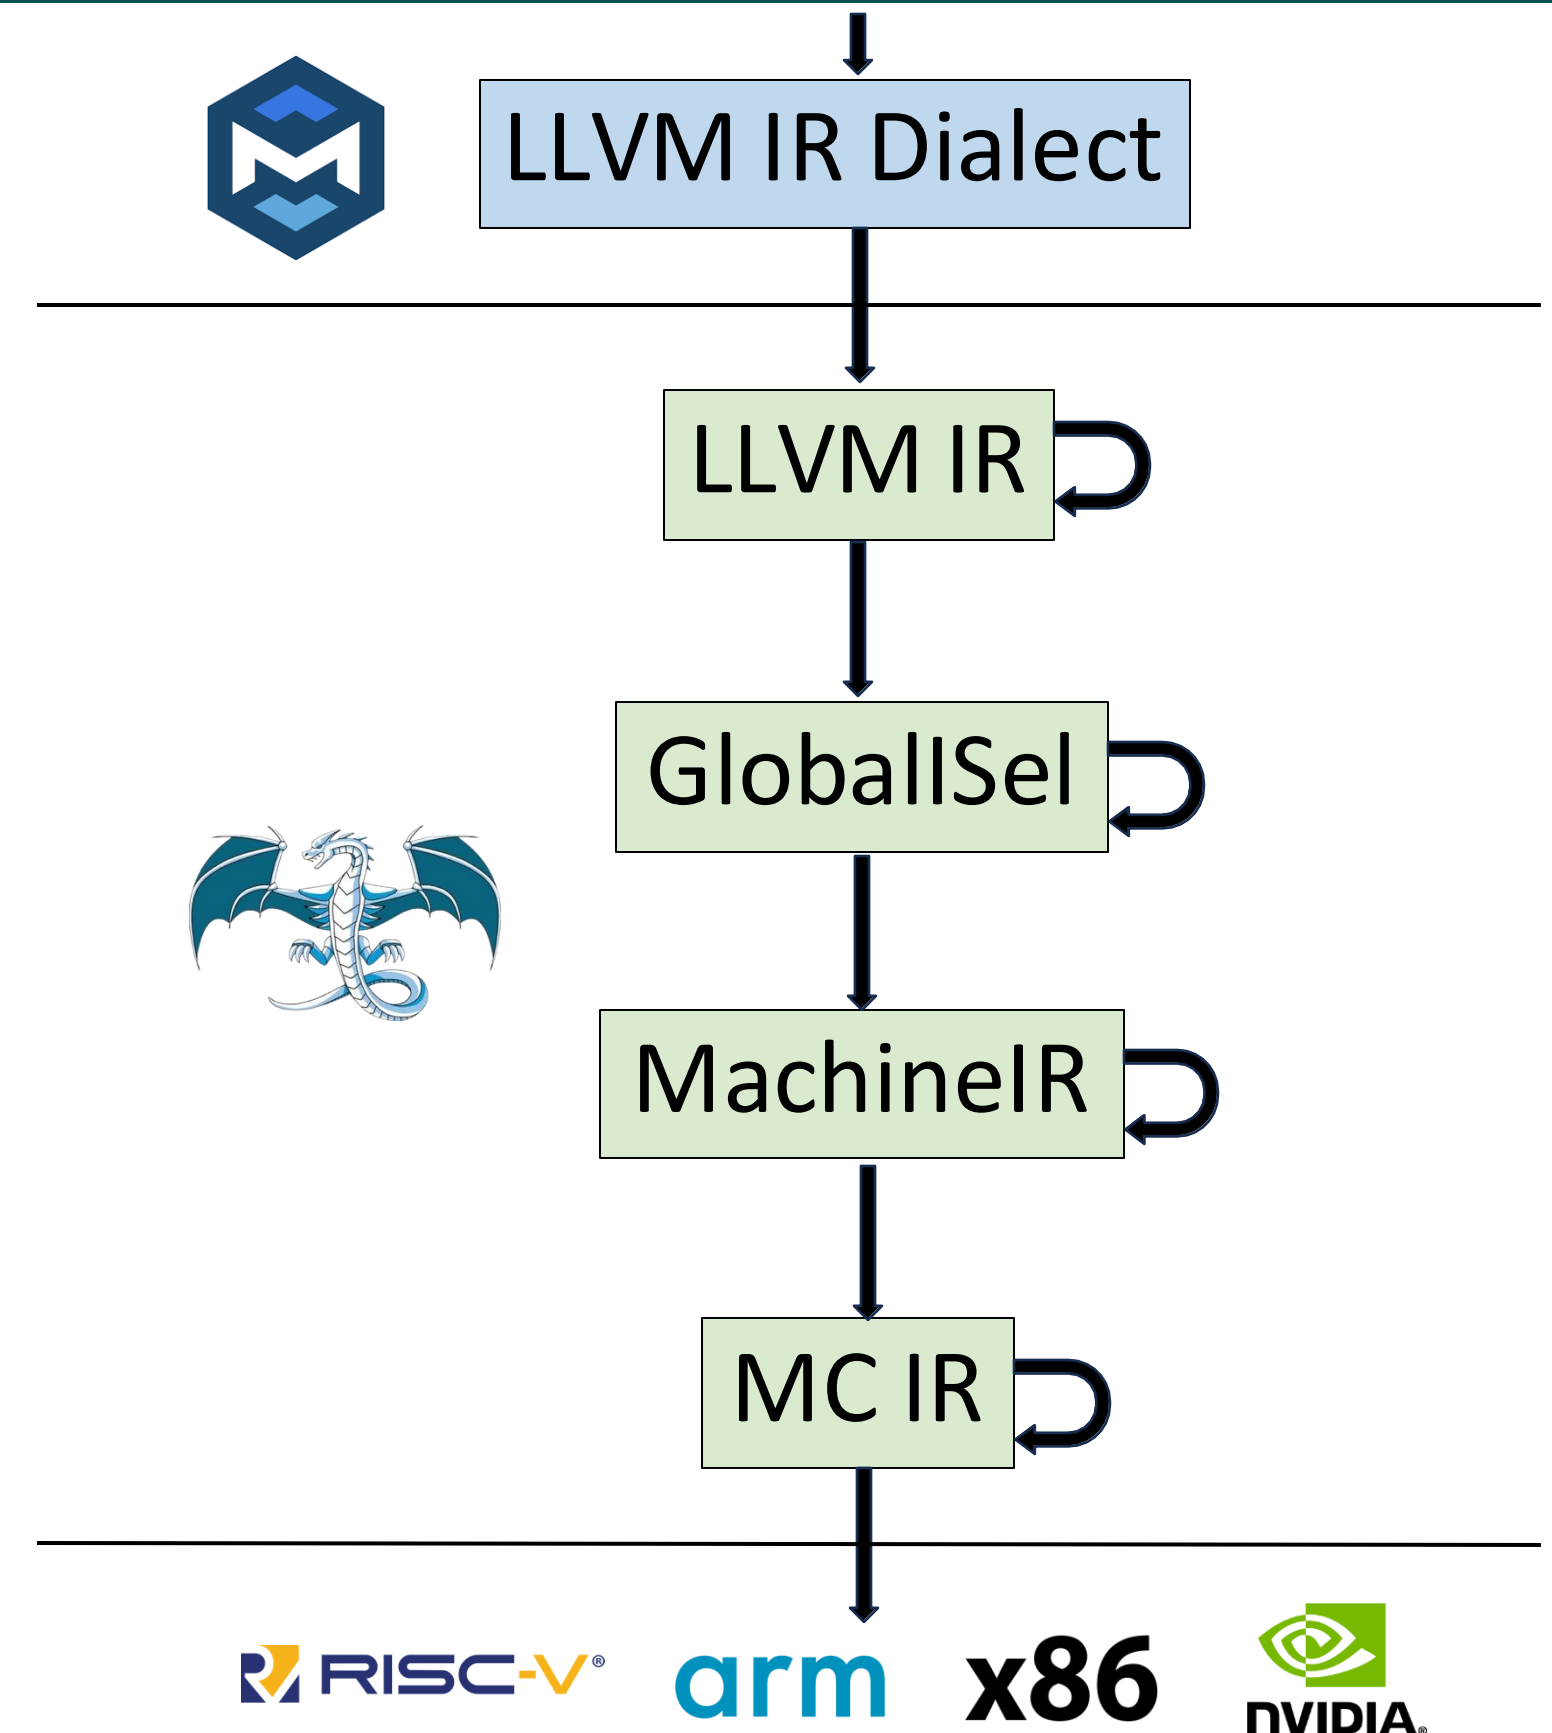
8/21/24



“大编译器” 设计 – LLVM后端



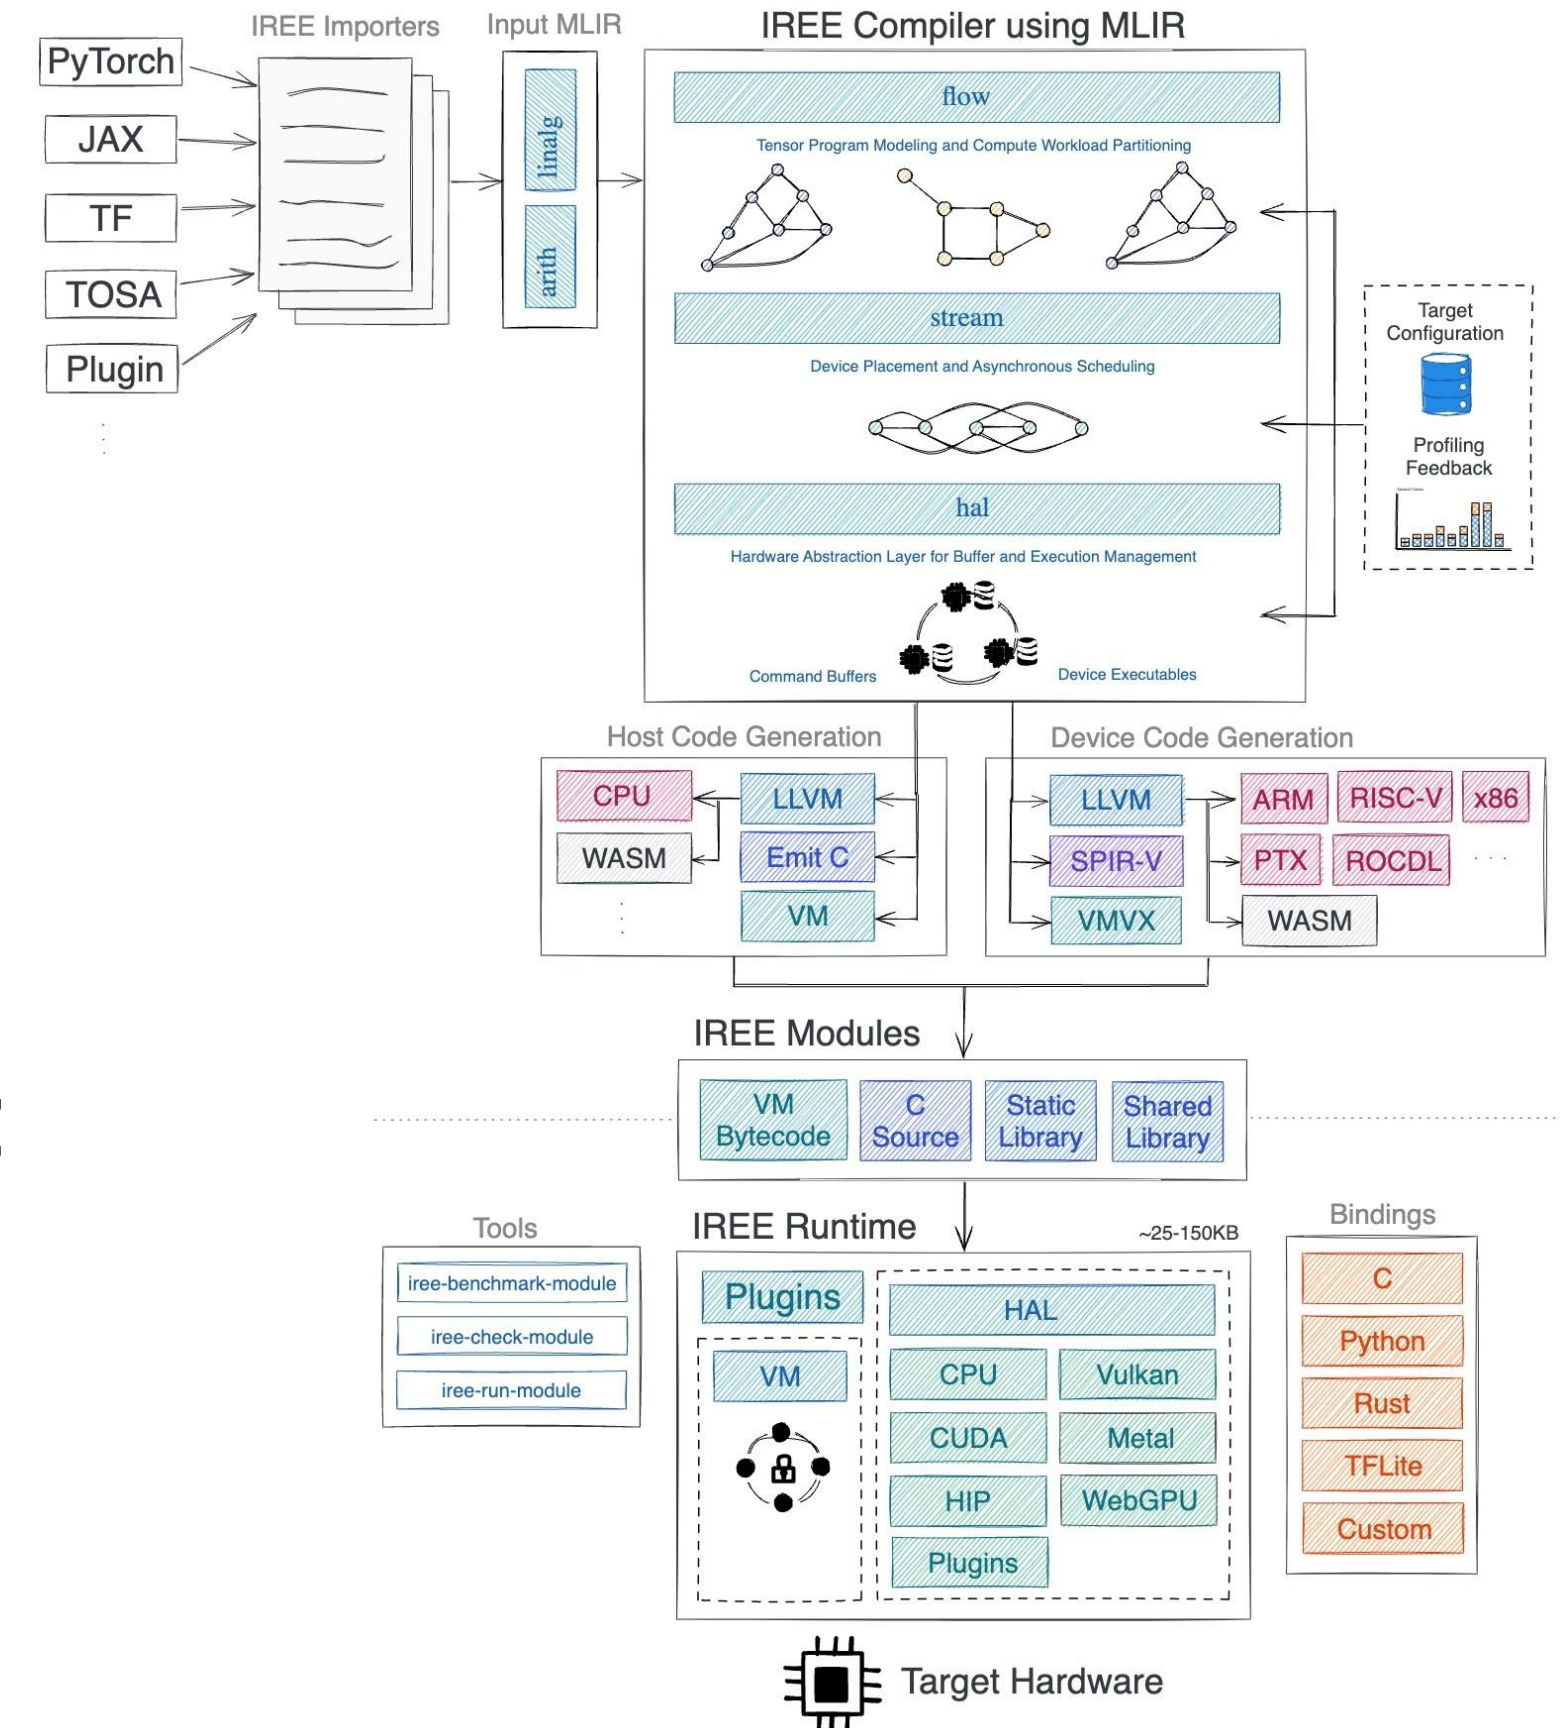
- 无缝衔接MLIR
- 支持多种硬件平台
- 处理标量相关的优化
- 指令选择，寄存器分配，指令排序等
- 快速适配各类RISC-V DSA



“大编译器” 实现 – 模型导入



- 基于IREE的模型导入
- 支持导入PyTorch, JAX/Tensorflow, ONNX
 - Dialect: TorchIR, StableHLO, TOSA
- 统一下降到LinalgExt, Arith, Flow etc



“大编译器” 实现 – 算子库



- Triton 算子库 -> MLIR Dialect TTIR
- 纯算法，跨平台的开源算子库
- MK (Magic Kernel) 抽象不同硬件的算子库
 - 通过MK映射到各硬件平台的算子库
 - 自动生成Kernel runtime

```
1 #any_device = #tt.operand_constraint<dram|l1|scalar|tile|any_device|any_device_tile>
2 module attributes {tt.system_desc = #tt.system_desc<[{arch = <wormhole_b0>, grid = <8x8>}]>
3 func.func @forward(%arg0: tensor<64x128xf32>, %arg1: tensor<64x128xf32>) -> tensor<64x128xf32>
4     %0 = tensor.empty() : tensor<64x128xf32>
5     %1 = "ttir.multiply"(%arg0, %arg1, %0) <{operandSegmentSizes = array<i32: 2, 1>, operands = [%arg0, %arg1, %0]}>
6     return %1 : tensor<64x128xf32>
7 }
8 }

5 module attributes {tt.device = #tt.device<tt.grid<8x8>, (d0, d1) -> (0, d0, d1)>, [0]>, {}}
6 func.func @forward(%arg0: tensor<64x128xf32, #layout>, %arg1: tensor<64x128xf32, #layout>) -> tensor<64x128xf32, #layout>
7     %0 = "ttmetal.alloc"() <{address = 262144 : i64, memory_space = #l1_, size = 32768 : i32}>
8     %1 = "ttmetal.host_write"(%arg0, %0) : (tensor<64x128xf32, #layout>, tensor<64x128xf32, #layout>) -> ()
9     %2 = "ttmetal.alloc"() <{address = 294912 : i64, memory_space = #l1_, size = 32768 : i32}>
10    %3 = "ttmetal.host_write"(%arg1, %2) : (tensor<64x128xf32, #layout>, tensor<64x128xf32, #layout>) -> ()
11    %4 = "ttmetal.alloc"() <{address = 327680 : i64, memory_space = #l1_, size = 32768 : i32}>
12    %5 = "ttmetal.dispatch"(%1, %3, %4) <{core_ranges = [#ttmetal.core_range<0x0, 1x1>, #ttmetal.core_range<0x1, 1x1>]}>
13    ^bb0(%arg2: !ttkernel.cb<0, 0, memref<64x128xf32, #l1_>>, %arg3: !ttkernel.cb<0, 1, memref<64x128xf32, #l1_>>) :
14        "ttkernel.cb_push_back"(%arg2) : (!ttkernel.cb<0, 0, memref<64x128xf32, #l1_>>) -> !ttkernel.cb<0, 0, memref<64x128xf32, #l1_>>
15        "ttkernel.return"() : () -> ()
16    }, {
17    ^bb0(%arg2: !ttkernel.cb<0, 0, memref<64x128xf32, #l1_>>, %arg3: !ttkernel.cb<0, 1, memref<64x128xf32, #l1_>>) :
18        "ttkernel.cb_push_back"(%arg3) : (!ttkernel.cb<0, 1, memref<64x128xf32, #l1_>>) -> !ttkernel.cb<0, 1, memref<64x128xf32, #l1_>>
19        "ttkernel.return"() : () -> ()
20    }, {
21    ^bb0(%arg2: !ttkernel.cb<0, 0, memref<64x128xf32, #l1_>>, %arg3: !ttkernel.cb<0, 1, memref<64x128xf32, #l1_>>) :
22        "ttkernel.builtin"(%arg2, %arg3, %arg4) <{kind = @eltwise, op = @multiply}> : (!ttkernel.cb<0, 0, memref<64x128xf32, #l1_>>, !ttkernel.cb<0, 1, memref<64x128xf32, #l1_>>) -> !ttkernel.cb<0, 0, memref<64x128xf32, #l1_>>
23        "ttkernel.return"() : () -> ()
24    }) : (tensor<64x128xf32, #layout1>, tensor<64x128xf32, #layout1>, tensor<64x128xf32, #layout1>) -> tensor<64x128xf32, #layout>
25    "ttmetal.dealloc"(%2) : (tensor<64x128xf32, #layout1>) -> ()
26    "ttmetal.dealloc"(%0) : (tensor<64x128xf32, #layout1>) -> ()
27    %6 = "ttmetal.alloc"() <{address = 0 : i64, memory_space = #system, size = 32768 : i32}>
28    %7 = "ttmetal.host_read"(%5, %6) : (tensor<64x128xf32, #layout1>, tensor<64x128xf32, #layout1>) -> tensor<64x128xf32, #layout>
29    "ttmetal.dealloc"(%4) : (tensor<64x128xf32, #layout1>) -> ()
30    return %7 : tensor<64x128xf32, #layout>
31 }
32 }
```


“大编译器” 实现 – 运行时抽象



- 和平台无关的运行时抽象
- 自动生成控制器代码
- 自动生成平台相关算子
- 算子库编写者无需考虑硬件平台API
 - 硬件平台API通过编译器隐藏和自动生成

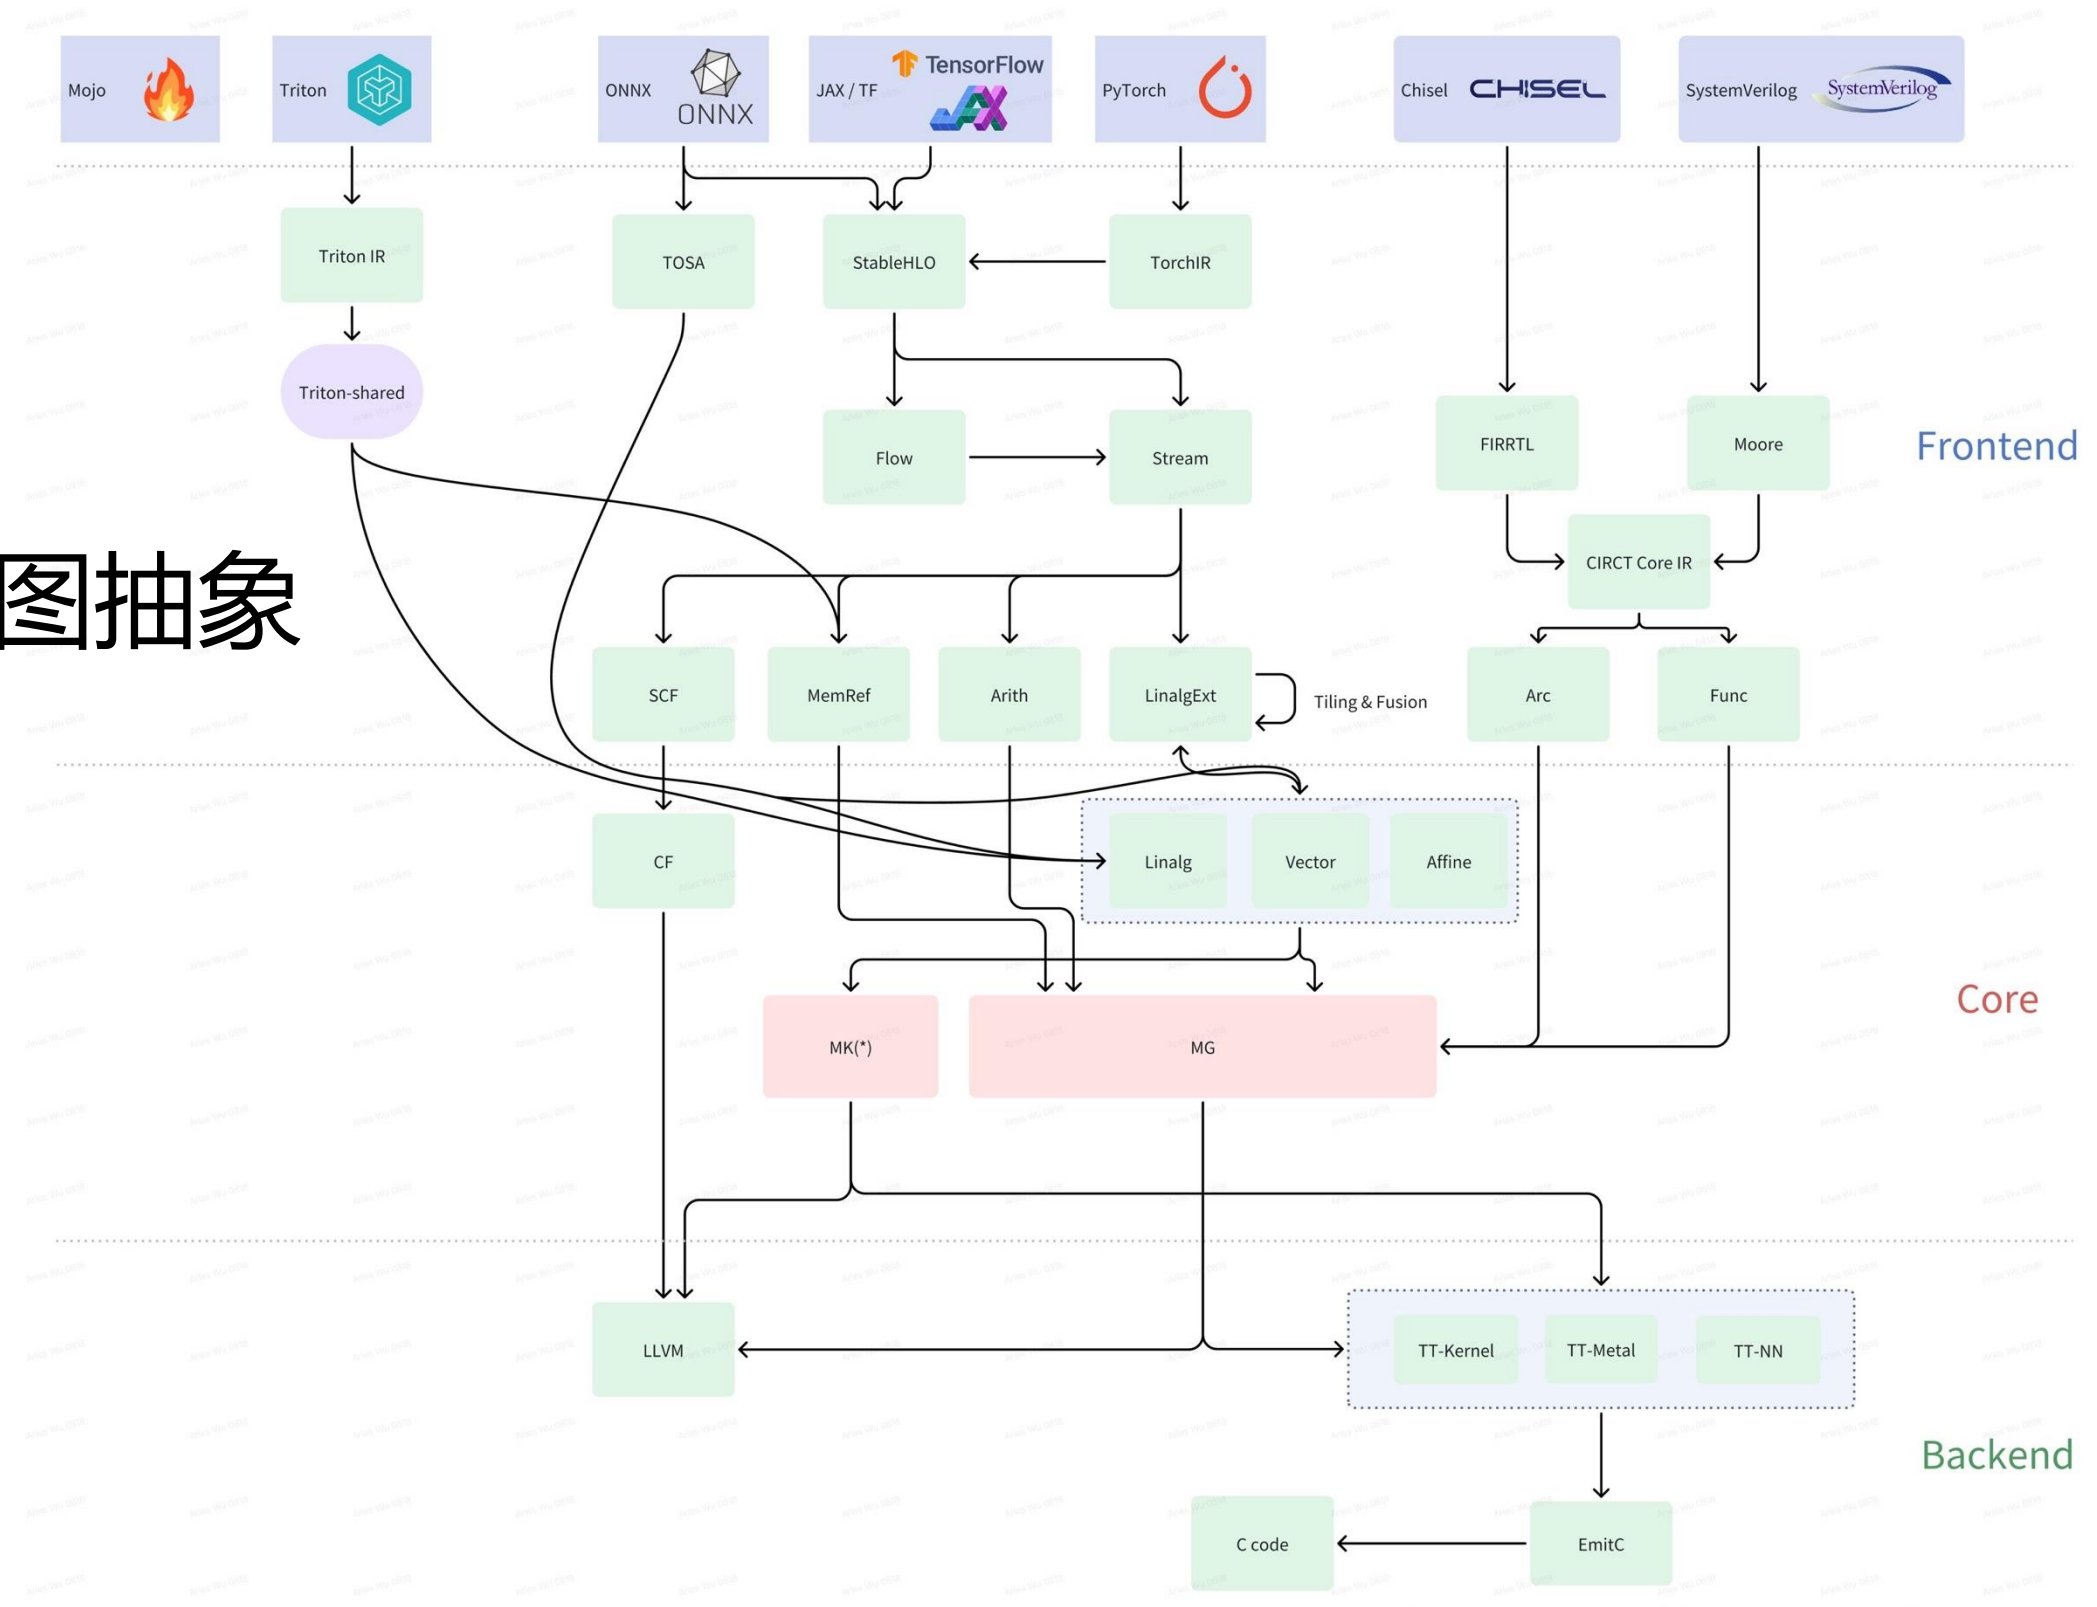
```
1 #any_device = #tt.operand_constraint<dram|l1|scalar|tile|any_device|any_device_tile>
2 module attributes {tt.system_desc = #tt.system_desc<[{arch = <wormhole_b0>, grid = <8x8>
3 func.func @forward(%arg0: tensor<64x128xf32>, %arg1: tensor<64x128xf32>) -> tensor<64x
4     %0 = tensor.empty() : tensor<64x128xf32>
5     %1 = "ttir.multiply"(%arg0, %arg1, %0) <{operandSegmentSizes = array<i32: 2, 1>, ope
6     return %1 : tensor<64x128xf32>|
7 }
8 }

5 module attributes {tt.device = #tt.device<#tt.grid<8x8>, (d0, d1) -> (0, d0, d1)>, [0]>,
6 func.func @forward(%arg0: tensor<64x128xf32, #layout>, %arg1: tensor<64x128xf32, #layo
7     %0 = "ttmetal.alloc"() <{address = 262144 : i64, memory_space = #l1_, size = 32768 :
8     %1 = "ttmetal.host_write"(%arg0, %0) : (tensor<64x128xf32, #layout>, tensor<64x128xf
9     %2 = "ttmetal.alloc"() <{address = 294912 : i64, memory_space = #l1_, size = 32768 :
10    %3 = "ttmetal.host_write"(%arg1, %2) : (tensor<64x128xf32, #layout>, tensor<64x128xf
11    %4 = "ttmetal.alloc"() <{address = 327680 : i64, memory_space = #l1_, size = 32768 :
12    %5 = "ttmetal.dispatch"(%1, %3, %4) <{core_ranges = [#ttmetal.core_range<0x0, 1x1>,
13    ^bb0(%arg2: !ttkernel.cb<0, 0, memref<64x128xf32, #l1_>>, %arg3: !ttkernel.cb<0, 1, i
14    "ttkernel.cb_push_back"(%arg2) : (!ttkernel.cb<0, 0, memref<64x128xf32, #l1_>>) ->
15    "ttkernel.return"() : () -> ()
16    }, {
17    ^bb0(%arg2: !ttkernel.cb<0, 0, memref<64x128xf32, #l1_>>, %arg3: !ttkernel.cb<0, 1, i
18    "ttkernel.cb_push_back"(%arg3) : (!ttkernel.cb<0, 1, memref<64x128xf32, #l1_>>) ->
19    "ttkernel.return"() : () -> ()
20    }, {
21    ^bb0(%arg2: !ttkernel.cb<0, 0, memref<64x128xf32, #l1_>>, %arg3: !ttkernel.cb<0, 1, i
22    "ttkernel.builtin"(%arg2, %arg3, %arg4) <{kind = @eltwise, op = @multiply}> : (!tt
23    "ttkernel.return"() : () -> ()
24    }) : (tensor<64x128xf32, #layout1>, tensor<64x128xf32, #layout1>, tensor<64x128xf32,
25    "ttmetal.dealloc"(%2) : (tensor<64x128xf32, #layout1>) -> ()
26    "ttmetal.dealloc"(%0) : (tensor<64x128xf32, #layout1>) -> ()
27    %6 = "ttmetal.alloc"() <{address = 0 : i64, memory_space = #system, size = 32768 : i
28    %7 = "ttmetal.host_read"(%5, %6) : (tensor<64x128xf32, #layout1>, tensor<64x128xf32,
29    "ttmetal.dealloc"(%4) : (tensor<64x128xf32, #layout1>) -> ()
30    return %7 : tensor<64x128xf32, #layout>
31 }
32 }
```

“大编译器” 实现 – 计算图



- Flow – AI模型计算图抽象
- Stream – 硬件流水线抽象
- MG (Magic Graph) – 通用计算图抽象
- 图切割算法
 - 动态计算时间模型
 - 数据通信模型
 - 全局变量重排
 - 缩减def-use链



Example - 简化的 GEMV in Triton



$Y = A @ X$, where A is a matrix of $M \times N$, X is a vector of N . @triton.jit

```
def gemv_kernel(Y, A, X, M, N, stride_am, BLOCK_SIZE_M: tl.constexpr, BLOCK_SIZE_N: tl.constexpr):
```

```
    start_m = tl.program_id(0)
```

```
    rm = start_m * BLOCK_SIZE_M + tl.arange(0, BLOCK_SIZE_M) rn =
```

```
    tl.arange(0, BLOCK_SIZE_N)
```

```
    A = A + (rm[:, None] * stride_am + rn[None, :])
```

```
    X = X + rn
```

```
    acc = tl.zeros((BLOCK_SIZE_M, ), dtype=tl.float32) for n in
```

```
    range(N, 0, -BLOCK_SIZE_N):
```

```
        a = tl.load(A)
```

```
        x = tl.load(X)
```


```
        acc += tl.sum(a * x[None, :], axis=1) A +=
```

```
        BLOCK_SIZE_N
```

```
        X += BLOCK_SIZE_N
```

```
    Y = Y + rm
```


```
    tl.store(Y, acc)
```

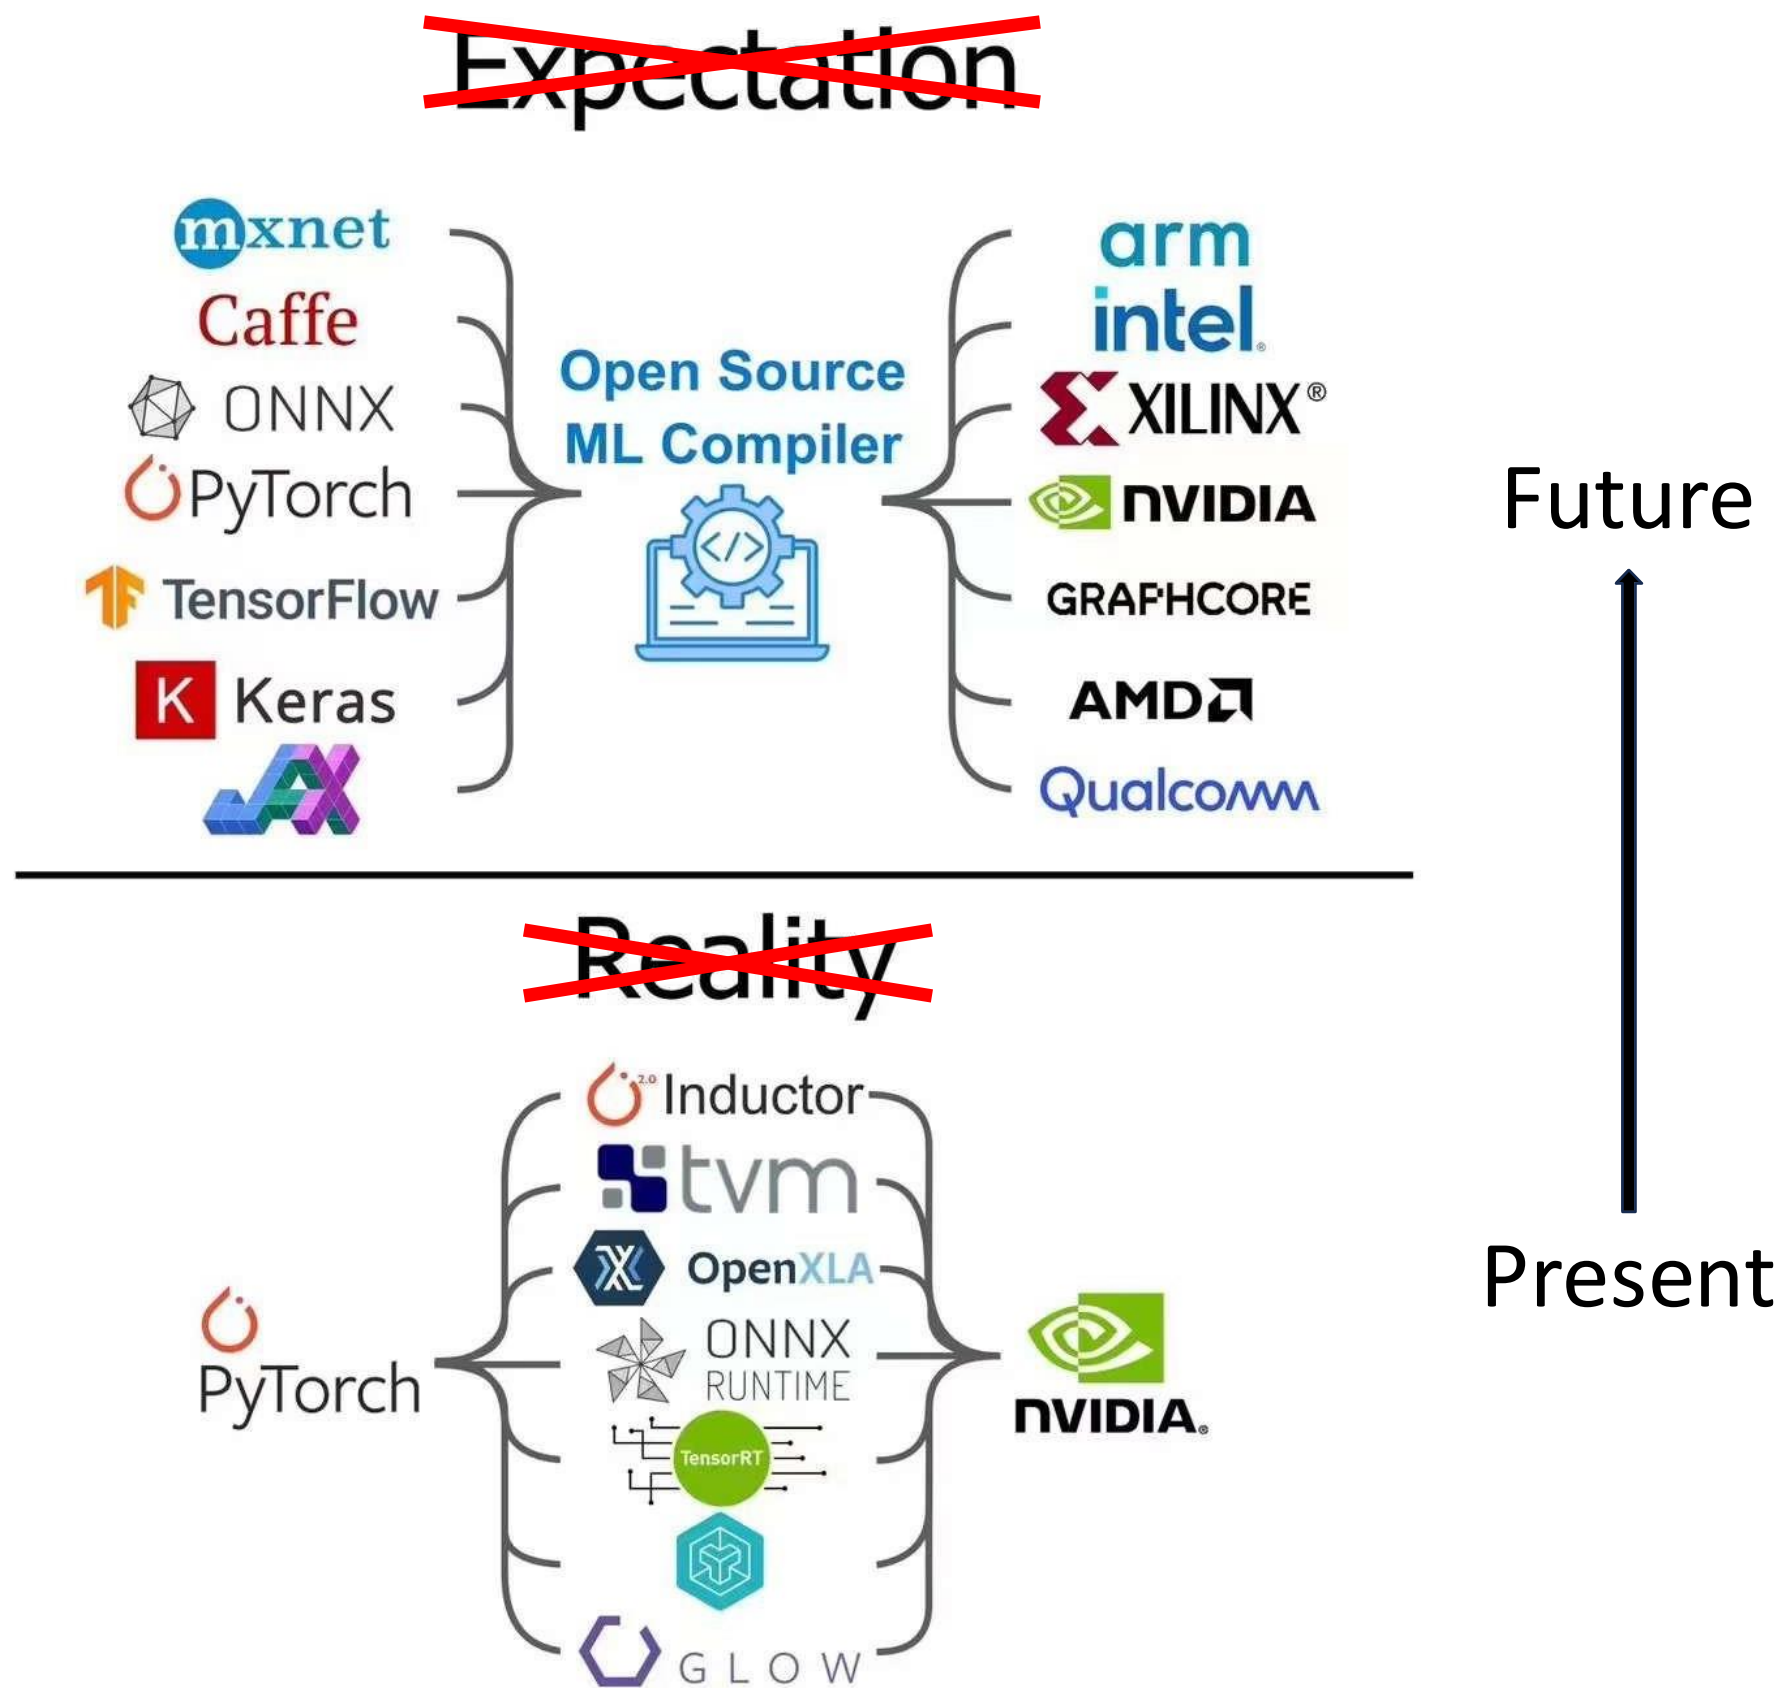


Example - 简化的 GEMV in TTIR



```
1 module {
2   tt.func public @gemv_kernel(%arg0: !tt.ptr<f32> loc("/path/to/gemv.py":10:0), %arg1: !tt.ptr<f32> loc("/path/to/gemv.py":10:0), %arg
3     %cst = arith.constant dense<0.000000e+00> : tensor<32xf32> loc(#loc1)
4     %c0_i32 = arith.constant 0 : i32 loc(#loc1)
5     %cst_0 = arith.constant dense<32> : tensor<32xi32> loc(#loc1)
6     %cst_1 = arith.constant dense<32> : tensor<32x32xi32> loc(#loc1)
7     %c32_i32 = arith.constant 32 : i32 loc(#loc1)
8     %0 = tt.get_program_id x : i32 loc(#loc2)
9     %1 = arith.muli %0, %c32_i32 : i32 loc(#loc3)
10    %2 = tt.make_range {end = 32 : i32, start = 0 : i32} : tensor<32xi32> loc(#loc4)
11    %3 = tt.splat %1 : i32 -> tensor<32xi32> loc(#loc5)
12    %4 = arith.addi %3, %2 : tensor<32xi32> loc(#loc5)
13    %5 = tt.expand_dims %4 {axis = 1 : i32} : tensor<32xi32> -> tensor<32x1xi32> loc(#loc6)
14    %6 = tt.splat %arg5 : i32 -> tensor<32x1xi32> loc(#loc7)
15    %7 = arith.muli %5, %6 : tensor<32x1xi32> loc(#loc7)
16    %8 = tt.expand_dims %2 {axis = 0 : i32} : tensor<32xi32> -> tensor<1x32xi32> loc(#loc8)
17    %9 = tt.broadcast %7 : tensor<32x1xi32> -> tensor<32x32xi32> loc(#loc9)
18    %10 = tt.broadcast %8 : tensor<1x32xi32> -> tensor<32x32xi32> loc(#loc9)
19    %11 = arith.addi %9, %10 : tensor<32x32xi32> loc(#loc9)
20    %12 = tt.splat %arg1 : !tt.ptr<f32> -> tensor<32x32x!tt.ptr<f32>> loc(#loc10)
21    %13 = tt.addptr %12, %11 : tensor<32x32x!tt.ptr<f32>>, tensor<32x32xi32> loc(#loc10)
22    %14 = tt.splat %arg2 : !tt.ptr<f32> -> tensor<32x!tt.ptr<f32>> loc(#loc11)
23    %15 = tt.addptr %14, %2 : tensor<32x!tt.ptr<f32>>, tensor<32xi32> loc(#loc11)
24    %16:3 = scf.for %arg6 = %c0_i32 to %arg4 step %c32_i32 iter_args(%arg7 = %cst, %arg8 = %13, %arg9 = %15) -> (tensor<32xf32>, tenc
25      %19 = tt.load %arg8 : tensor<32x32x!tt.ptr<f32>> loc(#loc13)
26      %20 = tt.load %arg9 : tensor<32x!tt.ptr<f32>> loc(#loc14)
27      %21 = tt.expand_dims %20 {axis = 0 : i32} : tensor<32xf32> -> tensor<1x32xf32> loc(#loc15)
28      %22 = tt.broadcast %21 : tensor<1x32xf32> -> tensor<32x32xf32> loc(#loc16)
29      %23 = arith.mulf %19, %22 : tensor<32x32xf32> loc(#loc16)
30      %24 = "tt.reduce"(%23) <{axis = 1 : i32}> ({
31        ^bb0(%arg10: f32 loc(unknown), %arg11: f32 loc(unknown)):
32          %28 = arith.addf %arg10, %arg11 : f32 loc(#loc29)
33          tt.reduce.return %28 : f32 loc(#loc27)
34      }) : (tensor<32x32xf32>) -> tensor<32xf32> loc(#loc27)
35      %25 = arith.addf %arg7, %24 : tensor<32xf32> loc(#loc20)
36      %26 = tt.addptr %arg8, %cst_1 : tensor<32x32x!tt.ptr<f32>>, tensor<32x32xi32> loc(#loc21)
37      %27 = tt.addptr %arg9, %cst_0 : tensor<32x!tt.ptr<f32>>, tensor<32xi32> loc(#loc22)
38      scf.yield %25, %26, %27 : tensor<32xf32>, tensor<32x32x!tt.ptr<f32>>, tensor<32x!tt.ptr<f32>>
39    } loc(#loc12)
40    %17 = tt.splat %arg0 : !tt.ptr<f32> -> tensor<32x!tt.ptr<f32>> loc(#loc24)
41    %18 = tt.addptr %17, %4 : tensor<32x!tt.ptr<f32>>, tensor<32xi32> loc(#loc24)
42    tt.store %18, %16#0 : tensor<32x!tt.ptr<f32>> loc(#loc25)
43    tt.return loc(#loc26)
44  } loc(#loc)
45 }
```





Thanks