

LLVM RISC-V Retrospect & Outlook

回顾与展望

肖玮 Wei Xiao (wei3.xiao@intel.com)

Contributors:

刘天乐 Tianle Liu (tianle.l.liu@intel.com)

王天晴 Tianqing Wang (tianqing.wang@intel.com)

闻浩海 Haohai Wen (haohai.wen@intel.com)

裴根 Gen Pei (gen.pei@intel.com)

Agenda

- LLVM RISC-V Target Support Overview
- LLVM Architecture Overview
- RISC-V & LLVM FE
- RISC-V & LLVM ME
- RISC-V & LLVM BE
- RISC-V Performance Analysis w/ QEMU

LLVM RISC-V Target Support Overview

▪ Base ISAs

Architecture	Description
riscv32	RISC-V with XLEN=32 (i.e. RV32I or RV32E)
riscv64	RISC-V with XLEN=64 (i.e. RV64I)

▪ Extensions

Table 109 Ratified Extensions by Status	
Extension	Status
A	Supported
B	Supported
C	Supported
D	Supported
F	Supported
E	Supported (See note)
H	Assembly Support
M	Supported
.....	

▪ Experimental Extensions

- ✓ experimental-zacas
- ✓ experimental-zihintntl
- ✓ experimental-zvfh
- ✓ experimental-zca
- ✓ ...

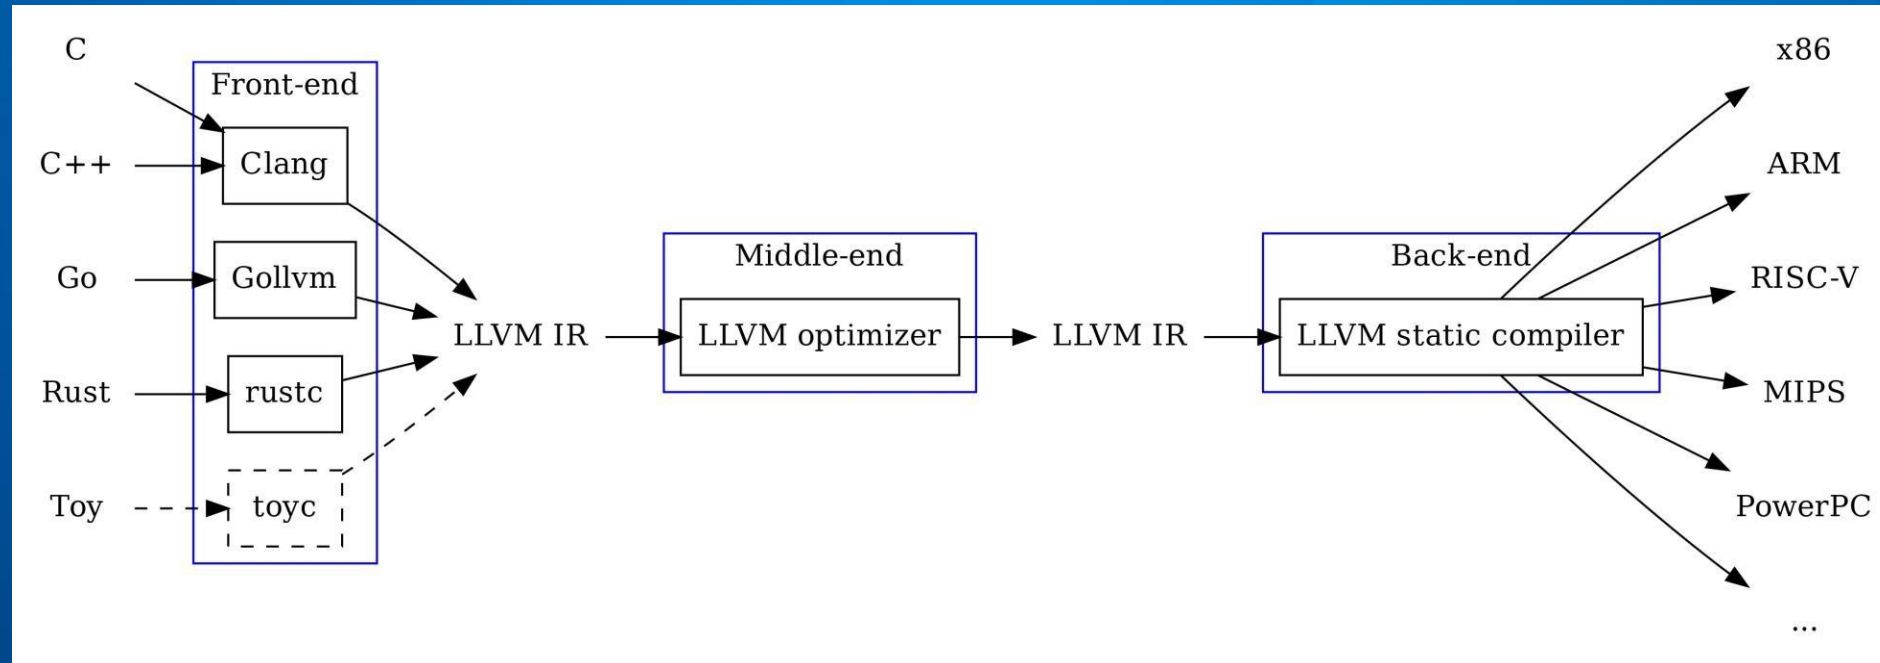
▪ Vendor Extensions

- ✓ XVentanaCondOps
- ✓ XtheadVdot
- ✓ XSfvcop
- ✓ ...

▪ Profiles

- ✓ rvi20u32
- ✓ rvi20u64
- ✓ ...

LLVM Architecture Overview



- Classical design to support multiple source languages or target architectures
- Serves a broader set of programmers which naturally leads to more enhancements and improvements to the compiler
- LLVM IR is a Complete Code Representation
- LLVM is a Collection of Libraries

RISC-V & LLVM FE

RISCV ABI generate **different** type for structure return type

Source code	X86 IR from FE	RISCV IR from FE	Comment of RISCV (Xlen =64)
<pre>struct A { int int }</pre>	<pre>%struct.A =type { i32, i32 } define dso_local i64 @_Z5myfunv()</pre>	<pre>%struct.A =type { i32, i32 } define dso_local i64 @_Z5myfunv()</pre>	Aggregates which are $\leq 2 \times \text{XLEN}$ will be passed in registers if possible
<pre>struct A { long int }</pre>	<pre>%struct.A =type { i64, i32 } define dso_local { i64, i32 } @_Z5myfunv()</pre>	<pre>%struct.A =type { i64, i32 } define dso_local [2 x i64] @_Z5myfunv()</pre>	Use a single XLEN int if possible, $2 \times \text{XLEN}$ if $2 \times \text{XLEN}$ alignment is required, and a 2-element XLEN array if only XLEN alignment is required.
<pre>struct A { long* int }</pre>	<pre>%struct.A =type { ptr, i32 } define dso_local { ptr, i32 } @_Z5myfunv()</pre>	<pre>%struct.A =type { ptr, i32 } define dso_local [2 x i64] @_Z5myfunv()</pre>	
<pre>struct A { long double }</pre>	<pre>%struct.A =type { i64, double } define dso_local { i64, double } @_Z5myfunv()</pre>	<pre>%struct.A =type { i64, double } define dso_local { i64, double } @_Z5myfunv()</pre>	Determine if a struct is eligible for passing according to the floating point calling convention (i.e., when flattened it contains a single fp value, fp+fp, or int+fp of appropriate size). If so, NeededArgFPRs and NeededArgGPRs are incremented appropriately.
<pre>struct A { long* int long };</pre>	<pre>%struct.A =type { ptr, i32, i64 } define dso_local void @_Z5myfunv(ptr noalias nocapture writeonly sret(%struct.A) align 8 %agg.result)</pre>	<pre>%struct.A =type { ptr, i32, i64 } define dso_local void @_Z5myfunv(ptr noalias nocapture writeonly sret(%struct.A) align 8 %agg.result)</pre>	

RISC-V & LLVM FE

"char" in RISC-V target

	X86-64 clang	Rv64 clang
bool	zeroext i1	zeroext i1
char	signext i8	zeroext i8
short	signext i16	signext i16
int	i32	signext i32
long	i64	i64

RISC-V

//i8_test.cpp:

```
char __attribute__((noinline)) foo() {
    return -1;
}

int main( int argc, char * argv[] )
{
    char ret;
    ret=foo();
    if( ret >=0)
        printf(">=0\n");
    else
        printf("<0\n");
}
```

Runtime Result

X86:

\$./i8_test.x86

<0

RISC-V:

\$./i8_test.rv

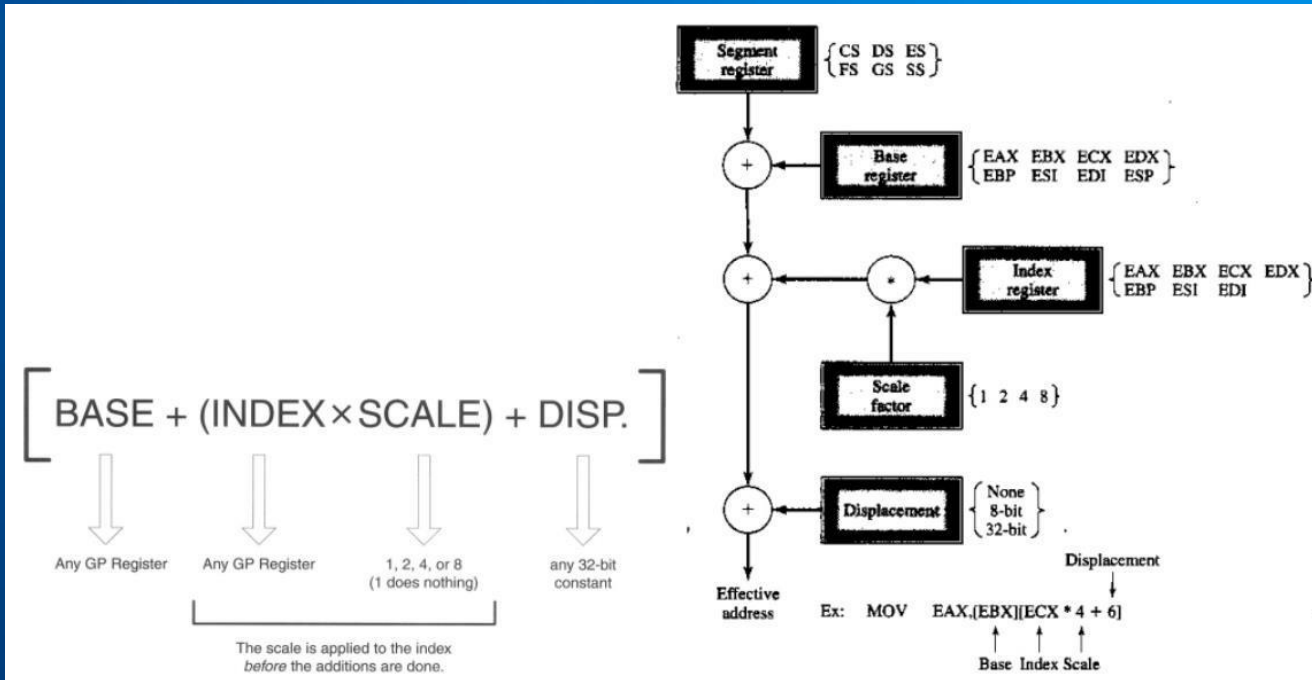
>=0

X86


```
386 ; Function Attrs: mustprogress norecurse uwtable
387 define dso_local noundef signext i32 @main(i32 noundef %argc, ptr noundef %argv) #1 {
388 entry:
389     %call = call noundef zeroext i8 @_Z3foov()
390     %conv = zext i8 %call to i32
391     %cmp = icmp sge i32 %conv, 0
392     br i1 %cmp, label %if.then, label %if.else
393
394 if.then:                                     ; preds = %entry
395     %call1 = call signext i32 (ptr, ...) @printf(ptr noundef @.str)
396     br label %if.end
397
398 if.else:                                     ; preds = %entry
399     %call2 = call signext i32 (ptr, ...) @printf(ptr noundef @.str.1)
400     br label %if.end
401
402 if.end:                                     ; preds = %if.else, %if.then
403     ret i32 0
404 }
```

```
386 ; Function Attrs: mustprogress norecurse uwtable
387 define dso_local noundef i32 @main(i32 noundef %argc, ptr noundef %argv) #1 {
388 entry:
389     %call = call noundef signext i8 @_Z3foov()
390     %conv = sext i8 %call to i32
391     %cmp = icmp sge i32 %conv, 0
392     br i1 %cmp, label %if.then, label %if.else
393
394 if.then:                                     ; preds = %entry
395     %call1 = call i32 (ptr, ...) @printf(ptr noundef @.str)
396     br label %if.end
397
398 if.else:                                     ; preds = %entry
399     %call2 = call i32 (ptr, ...) @printf(ptr noundef @.str.1)
400     br label %if.end
401
402 if.end:                                     ; preds = %if.else, %if.then
403     ret i32 0
404 }
```


RISC-V & LLVM ME: Addressing



X86



RISC-V

Register indirect with offset

```
ldr    x0, [Xn/sp, #imm]
ldr    x0, [Xn/sp]      ; #0 is implied if omitted
```

Register indirect with index

```
ldr    x0, [Xn/sp, Rn/zr, extend]
```

Extended	Effective address	Index format
[a, b, UXTW #0] [a, b, UXTW]	$a + (\text{uint32_t})b$	32-bit unsigned byte offset.
[a, b, UXTW #sizeshift]	$a + (\text{uint32_t})b * \text{size}$	32-bit unsigned element offset.
[a, b, SXTW #0] [a, b, SXTW]	$a + (\text{int32_t})b$	32-bit signed byte offset.
[a, b, SXTW #sizeshift]	$a + (\text{int32_t})b * \text{size}$	32-bit signed element offset.
[a, b, UCTX #0] [a, b, UCTX] [a, b, LSL #0] [a, b]	$a + (\text{uint64_t})b$	64-bit unsigned byte offset.
[a, b, UCTX #sizeshift] [a, b, LSL #sizeshift]	$a + (\text{uint64_t})b * \text{size}$	64-bit unsigned element offset.
[a, b, SCTX #0] [a, b, SCTX #sizeshift]	$a + (\text{int64_t})b$	64-bit signed byte offset.
	$a + (\text{int64_t})b * \text{size}$	64-bit signed element offset.

AArch64 (aka arm64)

RISC-V & LLVM ME: IPO

Inlining in RISC-V is less aggressive due to more cost

```
; cost before = -5, cost after = -5, .., cost delta = 0
  %arrayidx = getelementptr inbounds ptr, ptr %1, i64 %idxprom
...
Cost: 230
```

X86  Inlined

```
; cost before = -5, cost after = 0, .., cost delta = 5
  %arrayidx = getelementptr inbounds ptr, ptr %1, i64 %idxprom
...
Cost: 250
```

RISC-V  Not Inlined
(Default threshold: 250)

Different implementation of TargetLowering::isLegalAddressingMode() among the two targets

Outlook: more ME tuning works for RISC-V

RISC-V & LLVM ME: Vectorization

Scalable vectorization is enabled by default

```
for.body:
  %iv = phi i64 [ 0, %entry ], [ %iv.next, %for.body ]
  %arrayidx = getelementptr inbounds i64, ptr %a, i64 %iv
  %elem = load i64, ptr %arrayidx
  %divrem = udiv i64 %elem, %v
  store i64 %divrem, ptr %arrayidx
  %iv.next = add nuw nsw i64 %iv, 1
  %exitcond.not = icmp eq i64 %iv.next, 1024
  br i1 %exitcond.not, label %for.end, label %for.body
```

Loop Vectorizer

```
vector.ph:                                ; preds = %entry
  %2 = call i64 @llvm.vscale.i64()
  %3 = mul i64 %2, 2
  %n.mod.vf = urem i64 1024, %3
  %n.vec = sub i64 1024, %n.mod.vf
  %4 = call i64 @llvm.vscale.i64()
  %5 = mul i64 %4, 2
  %broadcast.splatinsert = insertelement <vscale x 2 x i64> poison, i64 %v, i64 0
  %broadcast.splat = shufflevector <vscale x 2 x i64> %broadcast.splatinsert, <vscale x 2 x i64>
  poison, <vscale x 2 x i32> zeroinitializer
  br label %vector.body

vector.body:                               ; preds = %vector.body, %vector.ph
  %index = phi i64 [ 0, %vector.ph ], [ %index.next, %vector.body ]
  %6 = add i64 %index, 0
  %7 = getelementptr inbounds i64, ptr %a, i64 %6
  %8 = getelementptr inbounds i64, ptr %7, i32 0
  %wide.load = load <vscale x 2 x i64>, ptr %8, align 8
  %9 = udiv <vscale x 2 x i64> %wide.load, %broadcast.splat
  store <vscale x 2 x i64> %9, ptr %8, align 8
  %index.next = add nuw i64 %index, %5
  %10 = icmp eq i64 %index.next, %n.vec
  br i1 %10, label %middle.block, label %vector.body, !llvm.loop !0

middle.block:                             ; preds = %vector.body
```

RISC-V & LLVM ME: Vectorization

Loop vectorization using fixed length vectors

```
for.body:
  %iv = phi i64 [ 0, %entry ], [ %iv.next, %for.body ]
  %arrayidx = getelementptr inbounds i64, ptr %a, i64 %iv
  %elem = load i64, ptr %arrayidx
  %divrem = udiv i64 %elem, %v
  store i64 %divrem, ptr %arrayidx
  %iv.next = add nuw nsw i64 %iv, 1
  %exitcond.not = icmp eq i64 %iv.next, 1024
  br i1 %exitcond.not, label %for.end, label %for.body
```

Loop Vectorizer

```
vector.ph:                                ; preds = %entry
  %broadcast.splatinsert = insertelement <4 x i64> poison, i64 %v, i64 0
  %broadcast.splat = shufflevector <4 x i64> %broadcast.splatinsert, <4 x i64> poison, <4 x i32>
  zeroinitializer
  br label %vector.body

vector.body:                               ; preds = %vector.body, %vector.ph
  %index = phi i64 [ 0, %vector.ph ], [ %index.next, %vector.body ]
  %0 = add i64 %index, 0
  %1 = add i64 %index, 4
  %2 = getelementptr inbounds i64, ptr %a, i64 %0
  %3 = getelementptr inbounds i64, ptr %a, i64 %1
  %4 = getelementptr inbounds i64, ptr %2, i32 0
  %5 = getelementptr inbounds i64, ptr %2, i32 4
  %wide.load = load <4 x i64>, ptr %4, align 8
  %wide.load1 = load <4 x i64>, ptr %5, align 8
  %6 = udiv <4 x i64> %wide.load, %broadcast.splat
  %7 = udiv <4 x i64> %wide.load1, %broadcast.splat
  store <4 x i64> %6, ptr %4, align 8
  store <4 x i64> %7, ptr %5, align 8
  %index.next = add nuw i64 %index, 8
  %8 = icmp eq i64 %index.next, 1024
  br i1 %8, label %middle.block, label %vector.body, !llvm.loop !0

middle.block:                              ; preds = %vector.body
```

RISC-V & LLVM BE

RISC-V "V" Vector Extension

- Scalable vectorization

```
.LBB0_2:                                # %vector.ph
    negw    a1, a2
    andi    a1, a1, 1024
    slli    a3, a3, 1
    mv      a4, a0
    mv      a5, a1
    vsetvli a6, zero, e64, m2, ta, ma
.LBB0_3:                                # %vector.body
    # =>This Inner Loop Header: Depth=1
    vl2re64.v    v8, (a4)
    vdivu.vx     v8, v8, s0
    vs2r.v      v8, (a4)
    sub         a5, a5, a2
    add         a4, a4, a3
    bnez        a5, .LBB0_3
# %bb.4:                                # %middle.block
    li         a2, 1024
    beq         a1, a2, .LBB0_7
.LBB0_5:                                # %scalar.ph
```

- Fixed length vectorization

```
    li      a2, 0
    lui     a3, 2
    add     a3, a0, a3
    vsetivli zero, 4, e64, m2, ta, ma
.LBB0_1:                                # %vector.body
    # =>This Inner Loop Header: Depth=1
    addi    a4, a0, 32
    vle64.v v8, (a0)
    vle64.v v10, (a4)
    vdivu.vx v8, v8, a1
    vdivu.vx v10, v10, a1
    vse64.v v8, (a0)
    vse64.v v10, (a4)
    addi    a0, a0, 64
    addi    a2, a2, 8
    bne     a0, a3, .LBB0_1
# %bb.2:                                # %for.end
```

RISC-V & LLVM BE Optimization Example

```
; 'zero' register not addressable in compressed store.
```

```
                =>  li a1, 0
sw zero, 0(a0)  =>  c.sw a1, 0(a0)
sw zero, 8(a0)  =>  c.sw a1, 8(a0)
sw zero, 4(a0)  =>  c.sw a1, 4(a0)
sw zero, 24(a0) =>  c.sw a1, 24(a0)
```

```
; compressed stores support limited offsets
```

```
lui a2, 983065  =>  lui a2, 983065
                =>  addi a3, a2, -256
sw  a1, -236(a2) =>  c.sw  a1, 20(a3)
sw  a1, -240(a2) =>  c.sw  a1, 16(a3)
sw  a1, -244(a2) =>  c.sw  a1, 12(a3)
sw  a1, -248(a2) =>  c.sw  a1, 8(a3)
sw  a1, -252(a2) =>  c.sw  a1, 4(a3)
sw  a0, -256(a2) =>  c.sw  a0, 0(a3)
```

RISCVMakeCompressible focused on reducing code size. It looks for cases where an instruction has been selected which can't be represented by one of the compressed (16-bit as opposed to 32-bit wide) instruction forms:

- register not being one of the registers addressable from the compressed instruction
- offset being out of range

RISC-V & LLVM BE Optimization Example

The Problem

- ♦ Middle-end optimizations replace `sext` instructions with `zext` if the sign bit is known to be zero.
- ♦ **i32→i64 `zext` is never cheaper than `sext` for RISC-V and `sext` can be free.**

```
void foo(int *x, int n) {  
    for (int i = 0; i < n; ++i)  
        x[i] += 1;  
}
```

RISC-V Assembly

```
foo:  
    blez    a1, .LBB0_2  
    slli    a1, a1, 32  
    srli    a1, a1, 32  
.LBB0_1:  
    ; loop body  
.LBB0_2:  
    ret
```

LLVM IR

```
define void @foo(ptr %x, i32 signext %n) {  
entry:  
    %cmp3 = icmp sgt i32 %n, 0  
    br i1 %cmp3, label %preheader, label %cleanup
```

```
preheader:  
    ; sign bit of i32 %n is known 0 here.  
    %wide.trip.count = zext i32 %n to i64  
    br label %for.body
```

```
for.body:  
    ...
```

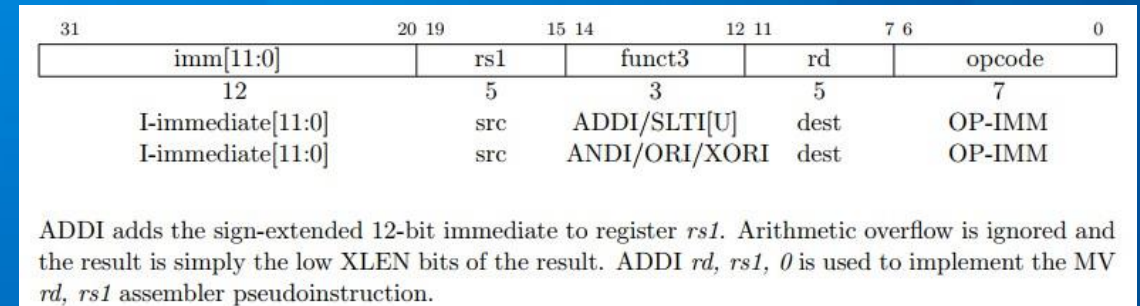
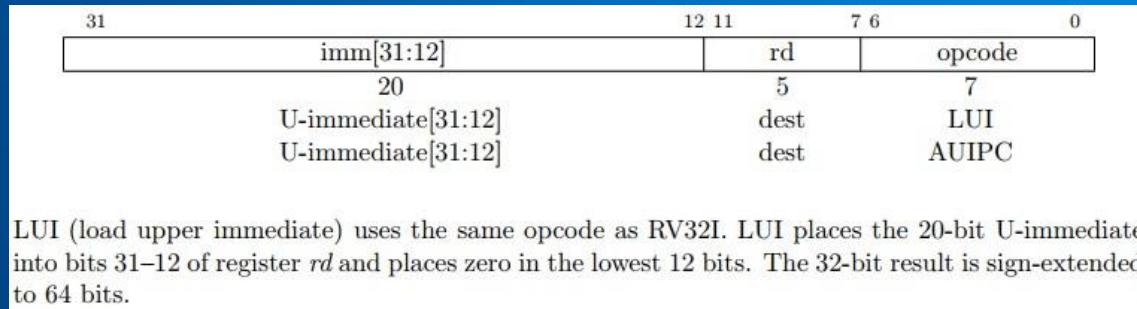
```
cleanup:  
    ret void
```

RISCVCodeGenPrepare

Partial Solution

- ♦ Pre-instruction selection IR pass
- ♦ For each `zext` instruction in basic blocks with a single predecessor
 - ◇ Examine the terminator condition of predecessor
 - ◇ If condition implies the sign bit is 0 when branching to the `zext`
 - ◇ Replace `zext` with `sext`
- ♦ Looking for a single predecessor is a very simple dominance check
 - ◇ **It will miss some cases.**

Macro-fusion LUI+ADDI(W) Optimization



a0=ptr @.str

```
lui    a0, %hi(.L.str)
fcvt.s.w    fa0, a1
addi    a0, a0, %lo(.L.str)
```

```
fcvt.s.w    fa0, a1
lui    a0, %hi(.L.str)
addi    a0, a0, %lo(.L.str)
```

uOP

3 → 2

a0=2048

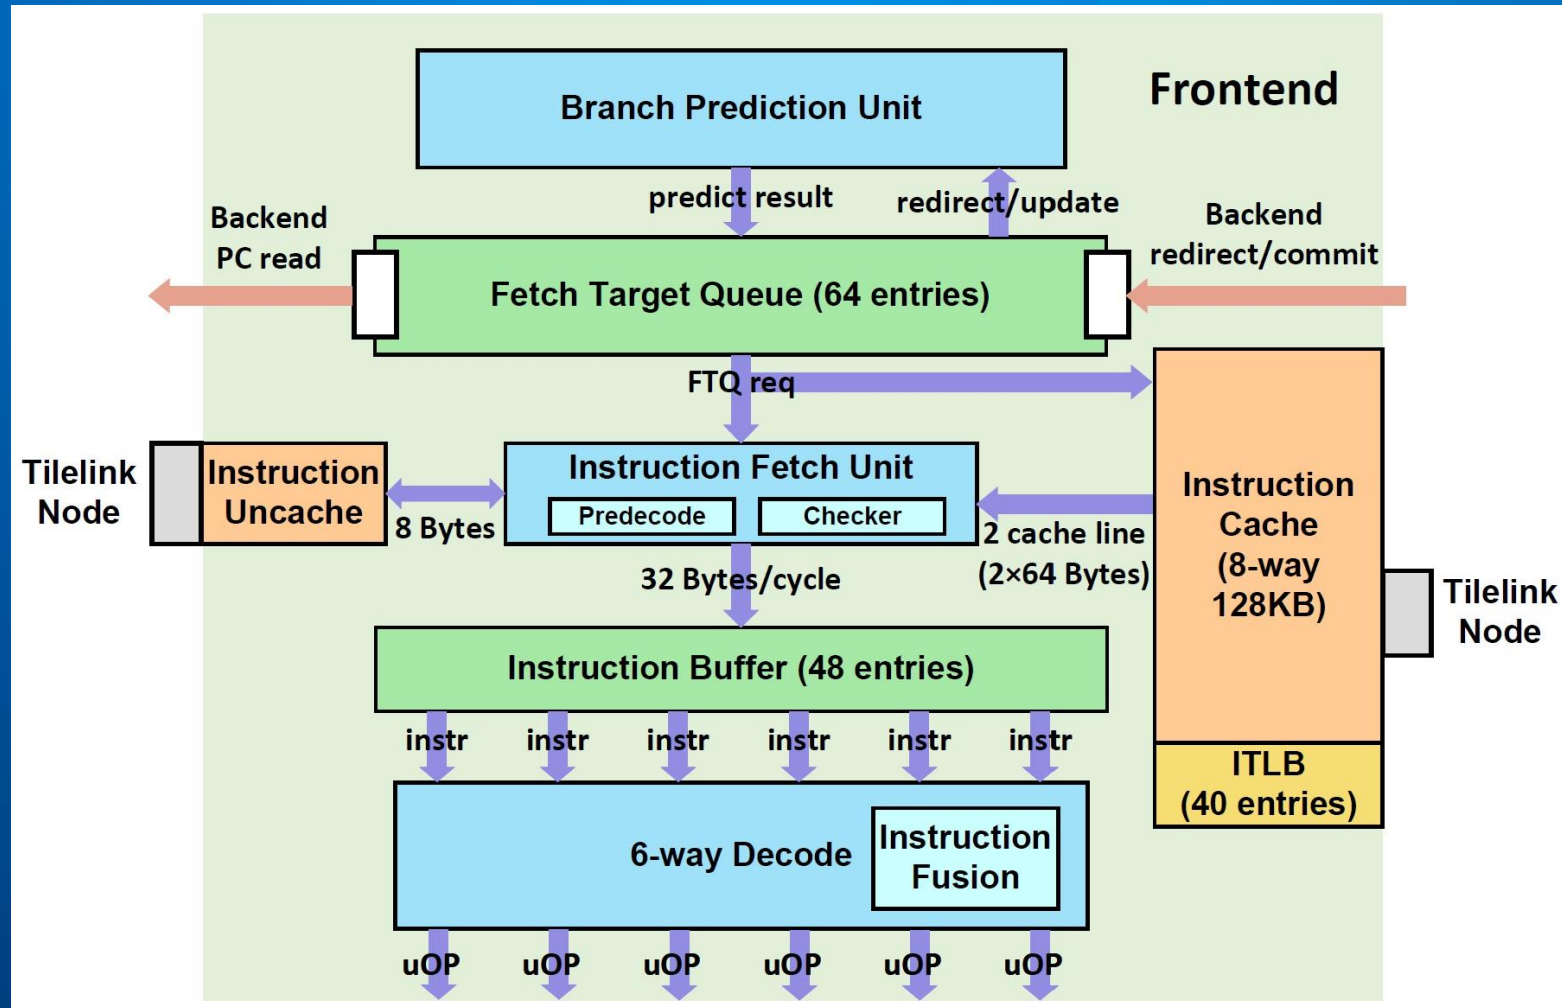
```
li      a0, 1
slli    a0, a0, 11
```

```
lui      a0, 1
addiw    a0, a0, -2048
```

2 → 1

Destination of LUI should be the ADDI(W) source register and destination register.

XiangShan Frontend



FusionDecoder & LLVM Optimizations

- clear upper 32 bits / get lower 32 bits: `slli r1,r0, 32 +srli r1,r1,32`
- clear upper 48 bits / get lower 16bits: `slli r1,r0, 48 +srli r1,r1,48`
- clear upper 48 bits / get lower 16bits: `slliw r1,r0, 16+srliw r1,r1,16`
- sign-extend a 16-bit number: `slliw r1,r0, 16+sraiw r1,r1,16`
- shift left by one and add: `slli r1,r0, 1+add r1,r1,r2`
- shift left by two and add: `slli r1,r0, 2 +add r1,r1,r2`
- shift left by three and add: `slli r1,r0, 3 +add r1,r1,r2`
- shift zero-extended word left by one: `slli r1,r0, 32 +srli r1,r0, 31`
- shift zero-extended word left by two: `slli r1,r0, 32 +srli r1,r0, 30`
- shift zero-extended word left by three: `slli r1,r0, 32 +srli r1,r0, 29`
- get the second byte: `srli r1,r0, 8 +andi r1,r1,255`
- shift left by four and add: `slli r1,r0, 4 +add r1,r1,r2`
- shift right by 29 and add: `srli r1,r0, 29 +add r1,r1,r2`
- shift right by 30 and add: `srli r1,r0, 30 +add r1,r1,r2`
- shift right by 31 and add: `srli r1,r0, 31 +add r1,r1,r2`
- shift right by 32 and add: `srli r1,r0, 32 +add r1,r1,r2`
- add one if odd, otherwise unchanged: `andi r1,r0, 1+add r1,r1,r2`
- add one if odd (in word format), otherwise unchanged: `andi r1,r0, 1+addw r1,r1,r2`
- addw and extract its lower 8 bits (fused into addwbyte)
- addw and extract its lower 1bit (fused into addwbit)
- addw and zext.h (fused into addwzext.h)
- addw and sext.h (fused into addwsext.h)
- logic operation and extract its LSB
- logic operation and extract its lower 16bits
- `OR(Cat(src1(63, 8), 0.U(8.W)), src2)`
- mul 7-bit data with 32-bit data

Outlook: more BE tuning works for RISC-V uArch

RISC-V Performance Analysis w/ QEMU

Linux perf tool limitations on some RISCV platforms:

1. Hard to get instruction-level perf data
2. Hard to distinguish between reading and writing for cache miss

Therefore, it is difficult to do in-depth analysis when analyzing data reading and writing performance bottlenecks.

RISC-V Performance Analysis w/ QEMU

Use QEMU to simulate cache behavior:

- `riscv-toolchain/qemu/contrib/plugins/cache.c`
- Compile `libmycache.so`
- Run `"qemu-riscv64 -cpu rv64 -plugin libmycache.so,dcachesize=n -d plugin ./app &> cache.log"`
- The default setting is 32 sets. Set according to the actual chip

Result Example

The cache miss information recorded by QEMU can be mapped to where the cache miss occurred in the bin. To carry out targeted optimization.

```
core #, data accesses, data misses, dmiss rate, insn accesses, insn misses, miss rate
0      112511961318 33690119436 29.9436% 448844998985 2697 0.0000%

address, data misses, instruction
0x166e8 (main), 1949979383, 00753027 fsc ft7,0(a0)
0x16730 (main), 1949979383, a110 fsc fa2,0(a0)
0x16720 (main), 1949978417, 0bc53c27 fsc ft8,184(a0)
0x16736 (main), 1949978401, a114 fsc fa3,0(a0)
0x1676e (main), 1949978376, 09856583 lwu a1,152(a0)
0x166f0 (main), 1949978342, 00653027 fsc ft6,0(a0)
0x16708 (main), 1949978330, 00053027 fsc ft0,0(a0)
0x1674a (main), 1949978180, 00353027 fsc ft3,0(a0)
0x16742 (main), 1949977669, 00553027 fsc ft5,0(a0)
0x16764 (main), 1949977639, a10c fsc fa1,0(a0)
0x166f8 (main), 1949977496, 01d53027 fsc ft9,0(a0)
0x16710 (main), 1949977478, 13053c27 fsc fa6,312(a0)
0x16752 (main), 1949977275, a118 fsc fa4,0(a0)
0x16714 (main), 1141819063, 07e53427 fsc ft10,104(a0)
0x16700 (main), 1141818937, 00253027 fsc ft2,0(a0)
0x16772 (main), 1064512913, 00053207 fld ft4,0(a0)
0x16718 (main), 986779842, 01153027 fsc fa7,0(a0)
0x16728 (main), 986741825, f8853027 fsc fs0,-128(a0)
0x1678a (main), 974989921, 2130 fld fa2,64(a0)
0x1677e (main), 974989490, 02053e07 fld ft8,32(a0)
0x16802 (main), 433807293, 37d2 fld fa5,304(sp)
0x16758 (main), 333657708, 00153027 fsc ft1,0(a0)
0x1673c (main), 178604113, a108 fsc fa0,0(a0)
0x16724 (main), 23509977, 00453027 fsc ft4,0(a0)
0x16958 (main), 15045491, fb87b007 fld ft0,-72(a5)
0x169ce (main), 14949837, 0307b987 fld fs3,48(a5)
0x1696c (main), 14949830, fd07b287 fld ft5,-48(a5)
0x1698c (main), 14949825, ff07b387 fld ft7,-16(a5)
0x169b0 (main), 14949761, 0107b0f7 fld ft9,16(a5)
```

```
166e2: 3432 fld fs0,296(sp)
166e4: 00c40533 add a0,a0,a2
166e8: 00753027 fsc ft7,0(a0) ; cache miss 1: 1,949,980,010
166ec: 00c48533 add a0,s1,a2
166f0: 00653027 fsc ft6,0(a0)
166f4: 00c78533 add a0,a5,a2
166f8: 01d53027 fsc ft9,0(a0)
166fc: 00c70533 add a0,a4,a2
16700: 00253027 fsc ft2,0(a0)
16704: 00c68533 add a0,a3,a2
16708: 00053027 fsc ft0,0(a0)
1670c: 00cf8533 add a0,t6,a2
16710: 13053c27 fsc fa6,312(a0)
16714: 07e53427 fsc ft10,104(a0)
16718: 01153027 fsc fa7,0(a0)
1671c: 00cd0533 add a0,s10,a2
16720: 00c53c27 fsc ft8,184(a0) ; cache miss 3: 1,949,978,417
16724: 00453027 fsc ft4,0(a0)
16728: f8853027 fsc fs0,-128(a0)
1672c: 00cf0533 add a0,t5,a2
16730: a110 fsc fa2,0(a0) ; cache miss 2: 1,949,979,383
16732: 00ce8533 add a0,t4,a2
16736: a114 fsc fa3,0(a0) ; cache miss 4: 1,949,978,401
16738: 00ce0533 add a0,t3,a2
1673c: a108 fsc fa0,0(a0)
1673e: 00c38533 add a0,t2,a2
16742: 00553027 fsc ft5,0(a0)
16746: 00c30533 add a0,t1,a2
1674a: 00353027 fsc ft3,0(a0)
1674e: 00c28533 add a0,t0,a2
16752: a118 fsc fa4,0(a0)
16754: 00c08533 add a0,a7,a2
```


Legal Disclaimer & Optimization Notice

- INFORMATION IN THIS DOCUMENT IS PROVIDED “AS IS”. NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO THIS INFORMATION INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.
- Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products.
- Copyright © 2024, Intel Corporation. All rights reserved. Intel, Pentium, Xeon, Xeon Phi, Core, VTune, Cilk, and the Intel logo are trademarks of Intel Corporation in the U.S. and other countries.

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804



intel®