

RISC-V Summit China 2024

# Enabling Hardware Sampling Based PGO for RISC-V Platform

Gao Yichuan ([yichuan.gao@intel.com](mailto:yichuan.gao@intel.com))

RISC-V Agile Design Lab

Intel Labs China



# PGO Basics



- PGO (Profile-Guided Optimization)
  - Use runtime **feedback** to improve software performance

## • Software vs Hardware PGO

### SWPGO (Instrumentation)

- Insert profiling **code snippet** to software
- Profiled data:
  - Branch/Jump count/target
  - Register/Memory values
- No hardware requirement
- High profiling overhead

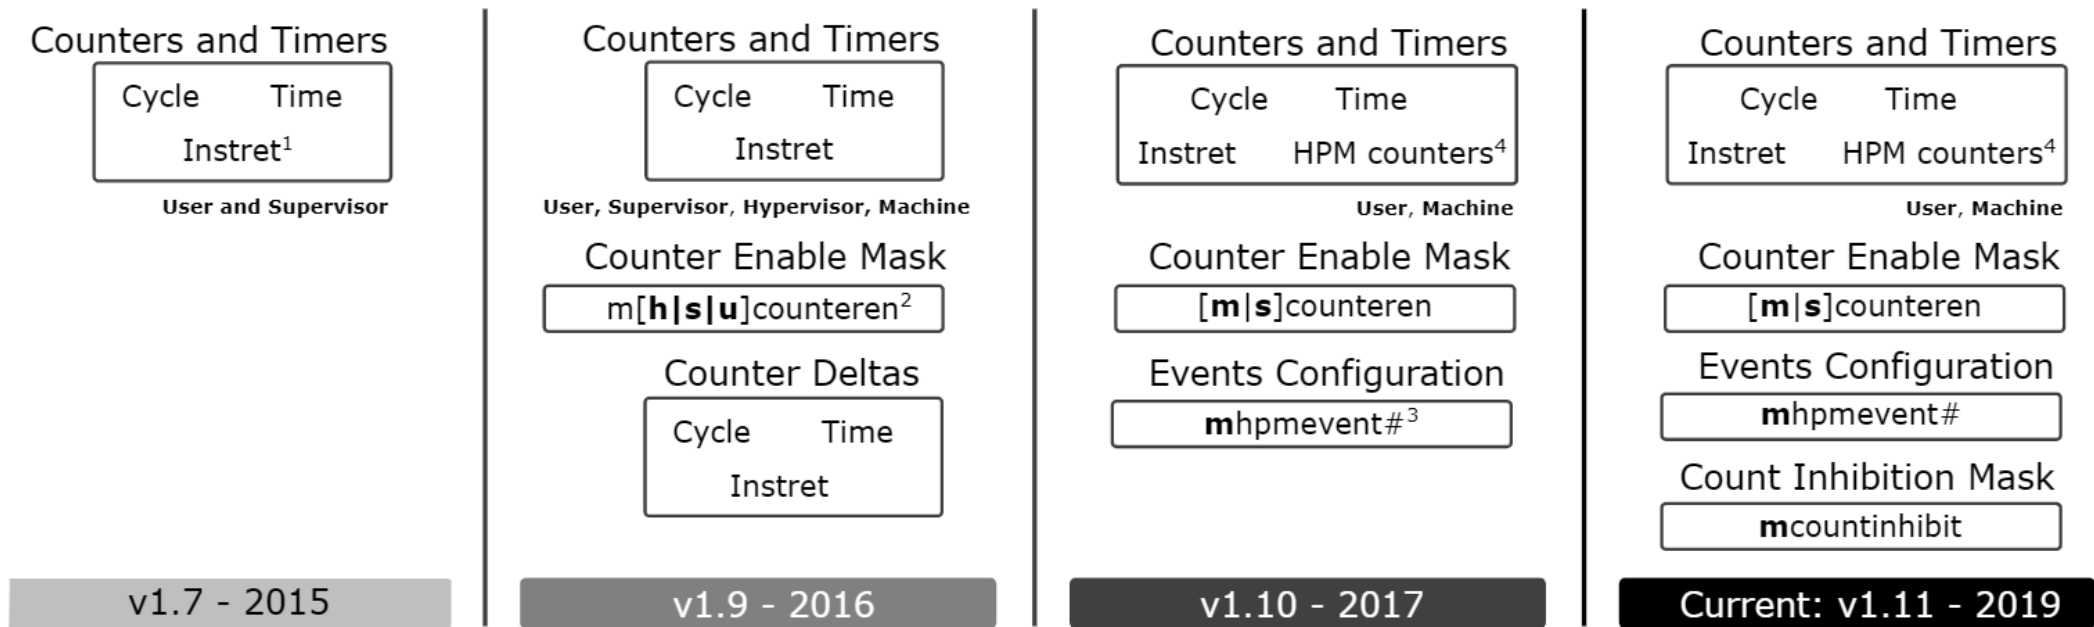
### HWPGO (Sampling)

- Use **hardware sample collect** methods
- Profiled data:
  - Performance Counter values
  - Branch/Jump count/target (with LBR/CTR)
  - Register/Memory values (not yet)
- Require **PMU hardware support**
- Low profiling overhead

This talk

# RISC-V PMU

- PMU (Performance Monitoring Unit)
  - Hardware unit for counting occurrence of **uArch events**
- RISC-V PMU: provide **fixed** and **event-based** hpmcounters
  - Software access via CSR interface



# RISC-V PMU Extensions for HWPGO

- Base extensions (widely adopted)
  - **Zicntr & Zihpm**  
Provide basic counters for PMU events

- Latest PMU extensions

- **Smcdeleg & Ssccfg**  
PMU counter delegation

PMU config in S-mode

- **Smcntrpmf & Sscofpmf**  
Counter mode filtering and overflow interrupt

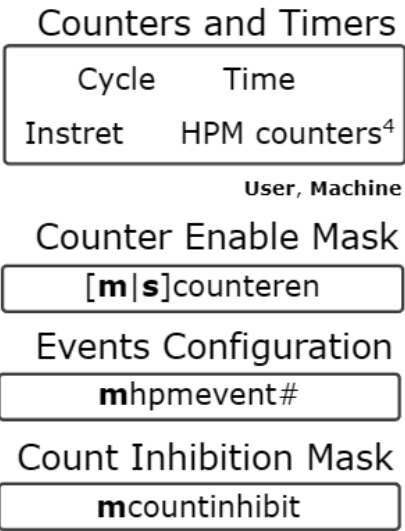
Event-based sampling

RVA23

- **Smctr & Ssctr**  
Hart control transfer records (in next slide)

Control flow info recording

- Future extensions for precise event sampling, register value profiling, etc.



# RISC-V CTR for HWPGO

- Record **control transfer history**
  - Jump instructions (including function calls and returns)
  - Taken/not-taken branch instructions
  - Traps, and trap returns
- Data is organized as circular buffer (FIFO)
  - Source PC, Target PC
  - Type, Cycle Count
  - Misprediction
- Support filtering by transfer type / privilege mode
- Low overhead sampling

## RISC-V Control Transfer Records (Smctr/Ssctr)

RISC-V Control Transfer Records TG

Version v1.0\_rc3, 2024-07-24: Frozen




Figure 5. Control Transfer Record Source Register Format for MXLEN=64




Figure 6. Control Transfer Record Target Register Format for MXLEN=64


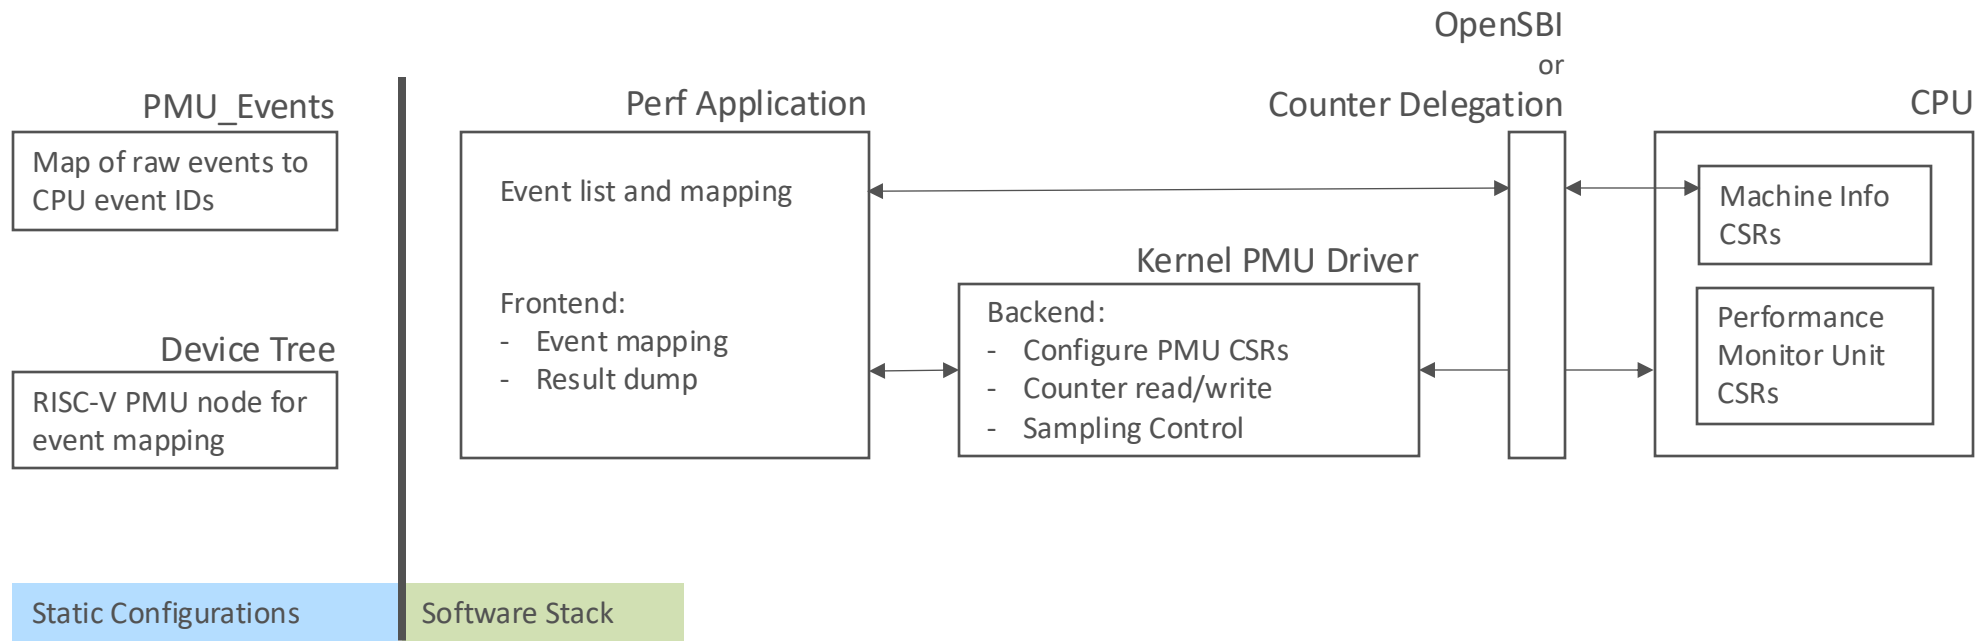


Figure 7. Control Transfer Record Metadata Register Format

Specifications: <https://github.com/riscv/riscv-control-transfer-records/>

# RISC-V Perf Software



# Case Study: Branch Prediction Profiling

- Perf tool sampling with CTR info

```
$ perf record -e branches:upp -e branch-misses:upp -b -z3 -- ./unpredictable
...
[ perf record: woken up 4033 times to write data ]
[ perf record: Captured and wrote 8.454 MB perf.data, compressed (original
1008.905 MB, ratio is 120.537) ]
```

```
$ perf report -i perf.data --header-only
# data offset      : 424
# data size       : 8864755
# feat offset     : 8865179
# hostname       : buildroot
# os release      : 6.9.0
# perf version    : 6.9.0
# arch           : riscv64
# cmdline        : /usr/bin/perf record -e branches:upp -e branch-misses:upp -b -z3 -- ./unpredictable
# event : name = branches:upp, , id = { 12 }, type = 0 (PERF_TYPE_HARDWARE), size = 136, config = 0x4
(PERF_COUNT_HW_BRANCH_INSTRUCTIONS), { sample_period, sample_freq } = 4000, sample_type = IP|TID|TIME|ID|PE>
# event : name = branch-misses:upp, , id = { 13 }, type = 0 (PERF_TYPE_HARDWARE), size = 136, config = 0x5
(PERF_COUNT_HW_BRANCH_MISSES), { sample_period, sample_freq } = 4000, sample_type = IP|TID|TIME|ID|PER>
# contains samples with branch stack
```

# Case Study: Mispredict Profile Feedback

```
1  for (int i = 0; i < N; i++) {
2      int *p;
3      if(s1[i] > 8000) {
4          p = &s2[i];
5          int z = i * i * i * i * i * i * i * i;
6          nop(p, z);
7      } else {
8          p = &s3[i];
9          nop(p, 3);
10     }
11     dst[i] = *p;
12 }
```

Poorly predicted from perf events:

- branches:upp
- branch-misses:upp

Attach 'unpredictable'  
Metadata to BR instruction

HWPGO

```
1  for (int i = 0; i < N; i++) {
2      int *p;
3
4      int c = (s1[i] > 8000);
5      p = c ? (&s2[i]) : (&s3[i]);
6      int z = i * i * i * i * i * i * i * i;
7      int arg2 = c ? z : 3;
8      nop(p, arg2);
9
10     dst[i] = *p;
11 }
```

Branch Predication

```
$ llvm-profgen --binary=unpredictable --output=unpredictable.prof --perfddata=perf.data --format=text
```

```
$ llvm-profgen --binary=unpredictable --output=unpredictable.misp.prof --perfddata=perf.data.4 --  
format=text --perf-event=branch-misses:upp --leading-ip-only
```

Mispredict profile generate (from perf.data) with llvm-profgen \*



# Case Study: Optimize with Clang

```
; if(s1[i] > 8000) {
105a0: 008a0533    add    a0, s4, s0
105a4: 4108        lw     a0, 0x0(a0)
105a6: 8b4e        mv     s6, s3
;    nop(p, z);
105a8: 00abc363    blt    s7, a0, 0x105ae <unpredictable+0x5a>
105ac: 8b4a        mv     s6, s2
;    nop(p, z);
105ae: fcabdbe3    bge    s7, a0, 0x10584 <unpredictable+0x30>
105b2: 02948533    mul    a0, s1, s1
105b6: 02950533    mul    a0, a0, s1
105ba: 029505b3    mul    a1, a0, s1
105be: 02a585bb    mulw   a1, a1, a0
```

Before HWPGO

Unpredictable branch instruction eliminated!

```
; if(s1[i] > 8000) {
10586: 008a0533    add    a0, s4, s0
1058a: 4108        lw     a0, 0x0(a0)
1058c: 00aba533    slt    a0, s7, a0
10590: 029485b3    mul    a1, s1, s1
10594: 029585b3    mul    a1, a1, s1
10598: 02958633    mul    a2, a1, s1
1059c: 02b605bb    mulw   a1, a2, a1
;    nop(p, z);
105a0: 0ea97633    czero.nez    a2, s2, a0
105a4: 0ea9d6b3    czero.eqz    a3, s3, a0
105a8: 8e55        or     a2, a2, a3
105aa: 00860b33    add    s6, a2, s0
105ae: 0ea5d5b3    czero.eqz    a1, a1, a0
105b2: 0eac7533    czero.nez    a0, s8, a0
105b6: 8dc9        or     a1, a1, a0
105b8: 855a        mv     a0, s6
105ba: f99ff0ef    jal    0x10552 <nop>
```

After HWPGO

Predication using RISC-V Zicond instruction

```
$ clang -march=rv64gc_zba_zbb_zbc_zbs_zicond -mabi=lp64d --target=riscv64-unknown-linux-gnu --
sysroot=riscv64-chroot -O3 -static -std=gnu99 nop.c unpredictable.c -fprofile-sample-
use=unpredictable.prof -mllvm -unpredictable-hints-file=unpredictable.misp.prof -o unpredictable_pgo
```

# Case Study: Performance Results

```
# perf stat -- ./unpredictable
```

```
Performance counter stats for './unpredictable':
```

```
178950.65 msec task-clock
          97      page-faults
2885889754      cycles
3858924390      instructions # 1.34  insn per cycle
```

```
# perf stat -- ./unpredictable_zicond
```

```
Performance counter stats for './unpredictable_zicond':
```

```
101149.07 msec task-clock
          97      page-faults
1630745749      cycles
4661493426      instructions # 2.86  insn per cycle
```

In misprediction demo application

(zicond with PGO)

**+21% Instructions**

**1.77x Performance**

**2.13x IPC**

In RISC-V Coremark (zicond extension w/o PGO)

**1.07x Performance**

**Call for community PGO work**

# Summary

- Hardware sampling-based PGO is beneficial for:
  - Lower overhead
  - Higher accuracy
  - New feedback capabilities for higher performance gains
- Call community collaboration on RISC-V PGO to:
  - Widely adopt latest PMU extensions (CTR and more) in hardware
  - Stabilize infrastructure for RISC-V Kernel Perf framework
  - Propose more PMU features for more feedback types
- Contact me via Email: [yichuan.gao@intel.com](mailto:yichuan.gao@intel.com)
- Special thanks to Intel SATG team for their invaluable work and support
  - More info: <https://llvm.org/devmtg/2024-04/slides/TechnicalTalks/Xiao-EnablingHW-BasedPGO.pdf>



# RISE

RISC-V Software Ecosystem

- <https://riseproject.dev>

**RISE is focused on positive and transparent collaborations with upstream projects to deliver commercial-ready software for various use cases**

**How:** Align on highest priorities & avoid (accidental) duplication of work

**Goal:** Accelerate open source SW for RISC-V architecture

<https://www.intel.com/content/www/us/en/developer/articles/community/rising-to-the-challenge-risc-v-software-readiness.html>

## Finding more interesting topics from Intel on RISC-V summit China 2024

Topic	When & Where
UXL 软件栈和 RISC-V 的初步探索	August 22 16:45 主会场A
LLVM 工具链 RISC-V 构建实现及其性能优化现状分析与未来展望	August 23 9:40 主会场A
GCC RVV 自动向量化及其应用	August 23 10:00 主会场A
Enhancing RISC-V Security with SBI Secure Service APIs	August 23 10:40 主会场B
Enabling Hardware Sampling Based PGO for RISC-V Platform	August 23 11:40 主会场A
利用 WASM 技术解决多种 ISA 的挑战	August 23 14:20 主会场B
HVP: Hardware Accelerated RISC-V Android Emulator	August 23 14:50 主会场A
Leverage BRS standard to improve RISC-V SW compatibility	August 23 17:30 主会场A
Soft-ISA: kernel built-in emulation engine to extend RISC-V silicon ISA capability	August 23 17:40 主会场A

# Legal Notices and Disclaimers

- Statements in this document that refer to future plans or expectations are forward-looking statements. These statements are based on current expectations and involve many risks and uncertainties that could cause actual results to differ materially from those expressed or implied in such statements. For more information on the factors that could cause actual results to differ materially, see our most recent earnings release and SEC filings at [www.intc.com](http://www.intc.com).
- All product plans and roadmaps are subject to change without notice. Any forecasts of goods and services needed for Intel's operations are provided for discussion purposes only. Intel will have no liability to make any purchase in connection with forecasts published in this document. Code names are often used by Intel to identify products, technologies, or services that are in development and usage may change over time. No license (express or implied, by estoppel or otherwise) to any intellectual property rights is granted by this document.
- © Intel Corporation. Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries. Other names and brands may be claimed as the property of others. This document contains information on products and/or processes in development.

The Intel logo is centered on a solid blue background. It features the word "intel" in a white, lowercase, sans-serif typeface. A small, bright blue square is positioned above the first vertical stroke of the letter 'i'. To the right of the word "intel" is a small white registered trademark symbol (®).

intel®